



# CS120: Computer Networks

## **Lecture 19. Other Topics in Transportation Layer**

Haoxian Chen


Slides adopted from: Zhice Yang

# Outline

- TCP Fairness
- QUIC
- QoS

# Evaluation Criteria

- Defining fairness is hard
  - In terms of a host, a TCP link, or an application ?
- TCP fairness goal: if **n** TCP sessions share same bottleneck link of bandwidth **R**, each should have average rate of ***R/n***



- Fairness Index



$$\bullet f(x_1 \dots x_n) = \frac{(\sum x_i)^2}{n * \sum x_i^2}$$

# Fairness in TCP

- Consider the steady state, TCP uses a (linear) scheme to adjust its window cwnd
  - $\text{cwnd}' = b * \text{cwnd} + a$
- Possible Designs
  - Additive increase, additive decrease
  - Additive increase, multiplicative decrease (AIMD)
  - Multiplicative increase, additive decrease
  - Multiplicative increase, multiplicative decrease


# Fairness in TCP

- Consider a case with two TCP connections



# Fairness in TCP


- Consider a case with two TCP connections



Increase the throughput to  
utilize more bandwidth


# Fairness in TCP

- Consider a case with two TCP connections



# Fairness in TCP


- Consider a case with two TCP connections






# Fairness in TCP

- Consider a case with two TCP connections




# Fairness in TCP

- Consider a case with two TCP connections
  - Behavior of additive increase additive decrease
    - Stable but not fair




# Fairness in TCP

- Consider a case with two TCP connections
  - Behavior of multiplicative increase multiplicative decrease
    - Stable but not fair




# Fairness in TCP

- Consider a case with two TCP connections
  - Behavior of multiplicative increase additive decrease
    - Not stable



# Fairness in TCP

- Consider a case with two TCP connections
  - Behavior of AIMD
    - Stable and fair



# Fairness and RTT

- TCP connection with smaller RTT occupies more bandwidth
  - When congestion happens, they recover more quickly
    - TCP adjust cwnd in RTT basis

# Fairness and Parallel TCP Connections

- Application can open multiple parallel connections between two hosts
  - web browsers do this , e.g., link of rate  $R$  with 9 existing connections:
    - new app asks for 1 TCP, gets rate  $R/10$
    - new app asks for 11 TCPs, gets  $R/2$

# Fairness and UDP

- Some apps do not use TCP
  - do not want rate throttled by congestion control
- Instead, use UDP:
  - send audio/video at constant rate, tolerate packet loss
- There is no “Internet police” policing use of congestion control



# Outline

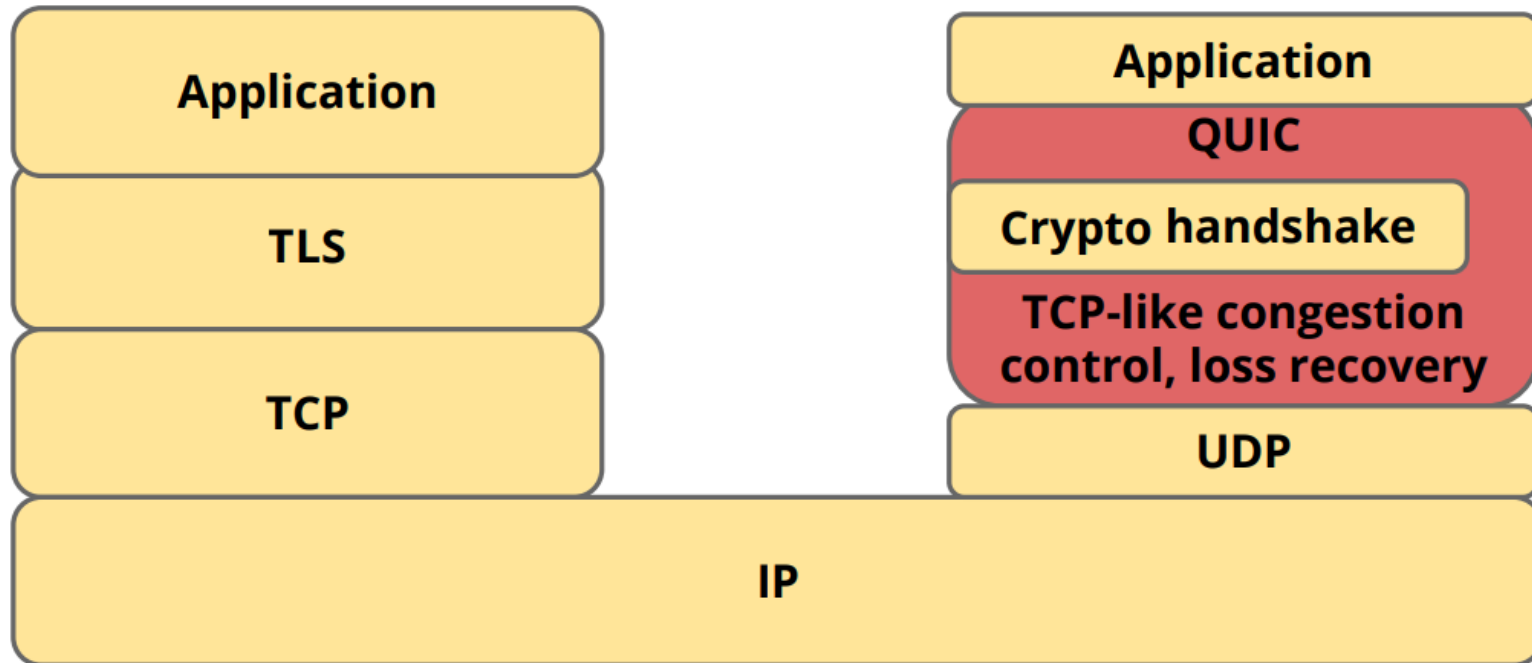
- TCP Fairness
- QUIC
- QoS

# QUIC

- QUIC: Quick UDP Internet Connections
- Application-layer protocol, on top of UDP
  - Deployed by Google starting at 2014
    - Deployed on many Google servers, apps (Chrome, mobile YouTube app)
  - QUIC working group formed in Oct 2016
- Initial goal: increase performance of HTTP

# QUIC

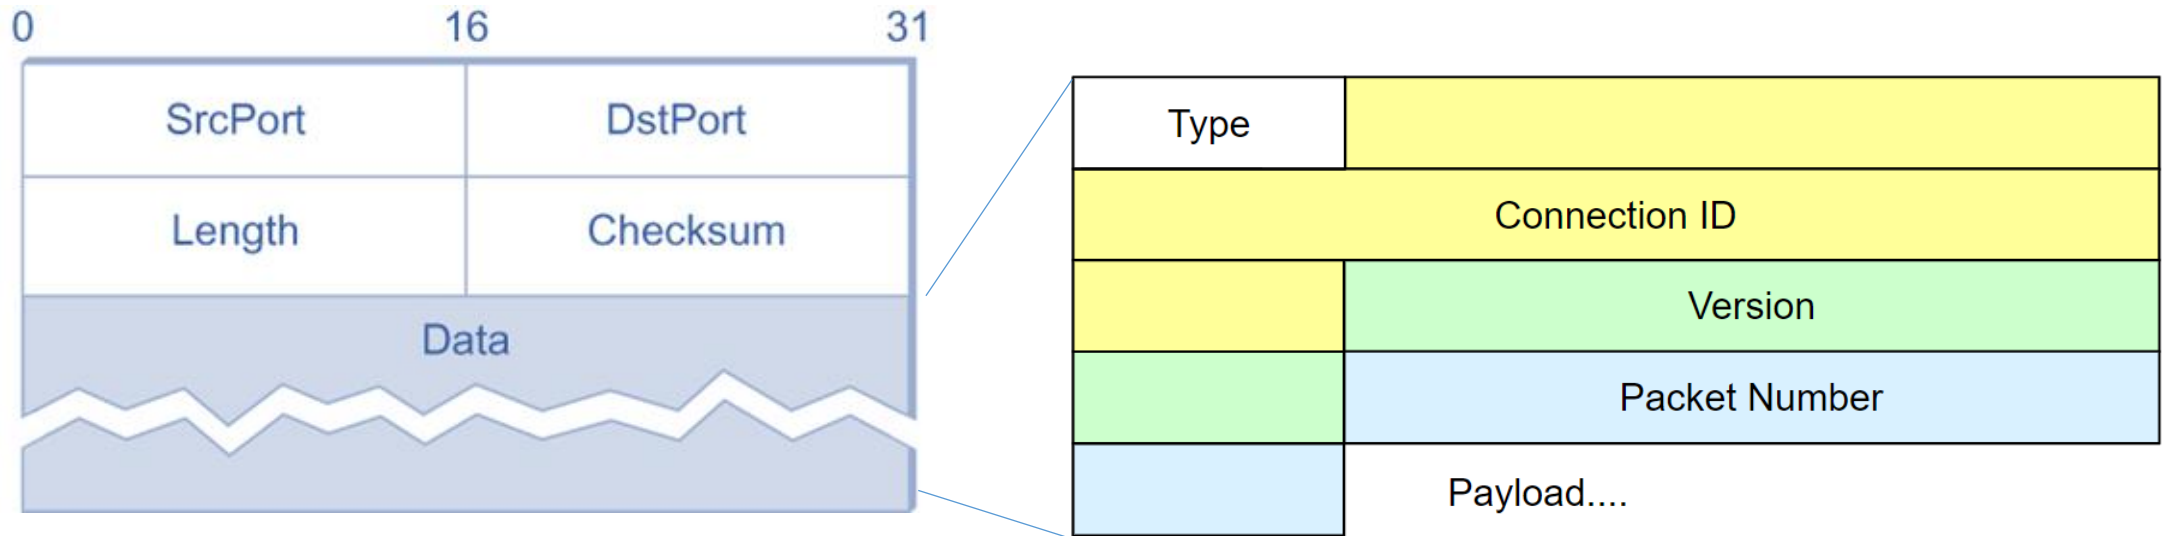
- Protocol Stack



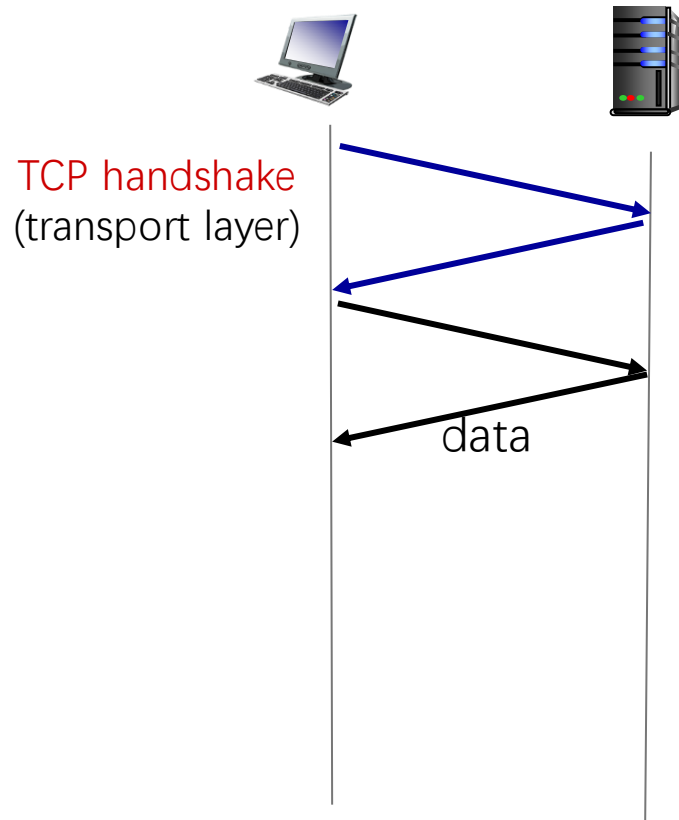
# QUIC

- Key features
  - Always encrypted
  - 0-RTT connection establishment
  - Connection migration
  - Congestion control
  - Parallel Streams

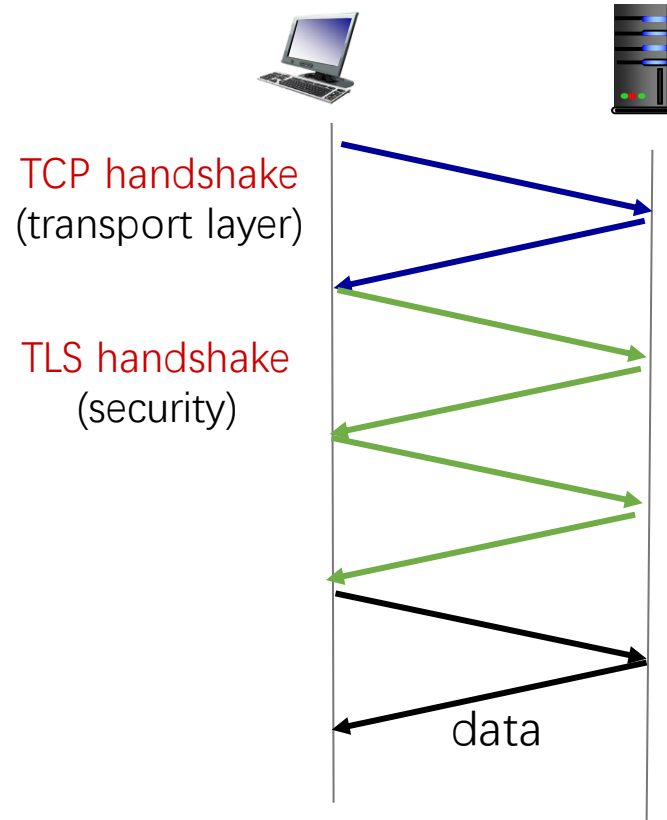
# QUIC - Header




# QUIC Connection Establishment



TCP  
(2RTT)



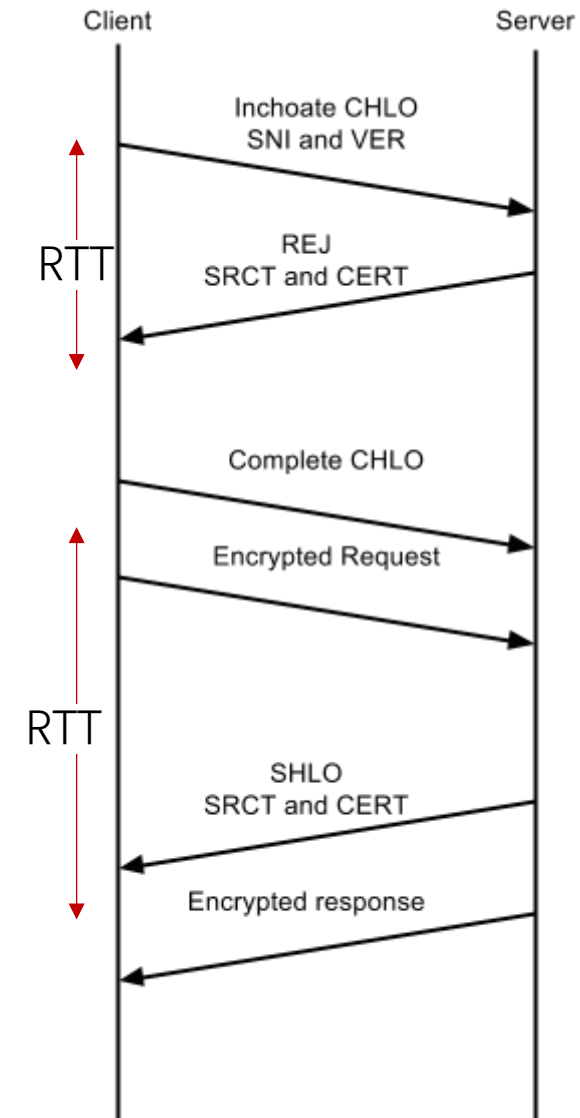
TCP+TLS 1.2  
(new 4RTT  
resumed 3RTT )



QUIC  
(new 2RTT  
Resumed 1RTT)

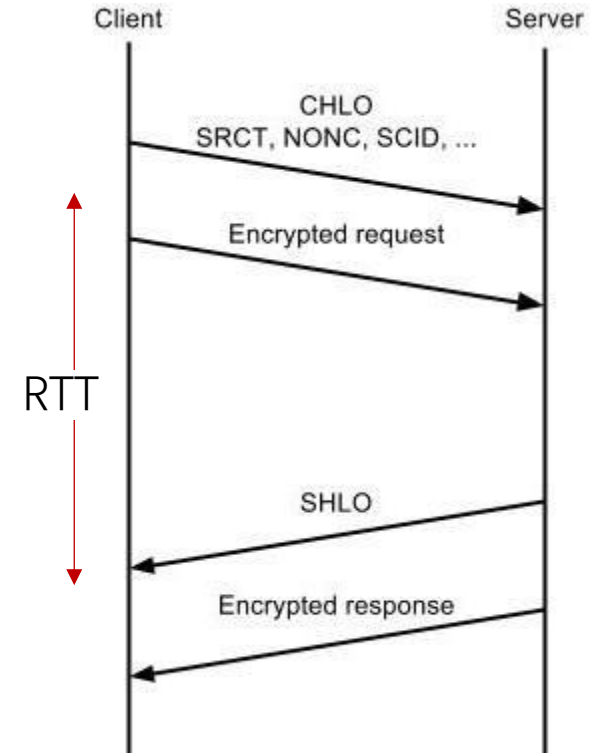
# QUIC Connection Establishment

- 1-RTT (First-ever connection)
  - No cached information available
  - First CHLO is inchoate (empty)
    - Simply includes version and server name
  - Server responds with REJ
    - Includes server config, certs, etc.
    - Allows client to make forward progress
  - Second CHLO is complete
    - Followed by initially encrypted request data
  - Server responds with SHLO
    - Followed immediately by forward-secure encrypted response data



# QUIC Connection Establishment

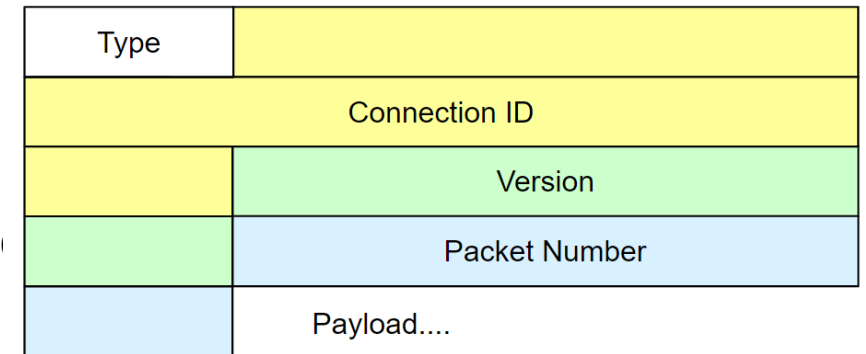
- 0-RTT (Subsequent connection)
  - Motivation: client can cache information about the *origin* it connected to
  - First CHLO is complete
    - Based on information from previous connection
    - Followed by initially encrypted data.
  - Server responds with SHLO
    - Followed immediately by forward-secure encrypted data





# QUIC Connection Migration

- NAT Rebinding
  - NATs remaps port
    - Frequency (~ mins)
    - Why ? to release unused ports
      - According to TCP connection state (if they are closed)
    - UDP does not have connection state, QUIC state is encrypted
- Mobility
  - Switching between different IP
    - Wi-Fi and cellular network
- Connection Migration
  - Keep QUIC connections alive even if port and IP are changing
  - Detect connection path changes via Connection ID and IP/port
    - Connection is identified by connection ID rather than <IP, port>
    - 64-bit connection ID
    - randomly chosen by client



# QUIC Congestion Control

- Incorporates TCP best practices
  - TCP Cubic, Fast Retransmission, Selective ACK, etc.
- Better signaling than TCP
  - Each packet carries a monotonically increasing packet number
    - Better RTT measurement
  - Retransmitted packets also consume new sequence numbers
    - no retransmission ambiguity
- More verbose ACK
  - support 256 ACK ranges (vs. TCP's 3 SACK ranges)

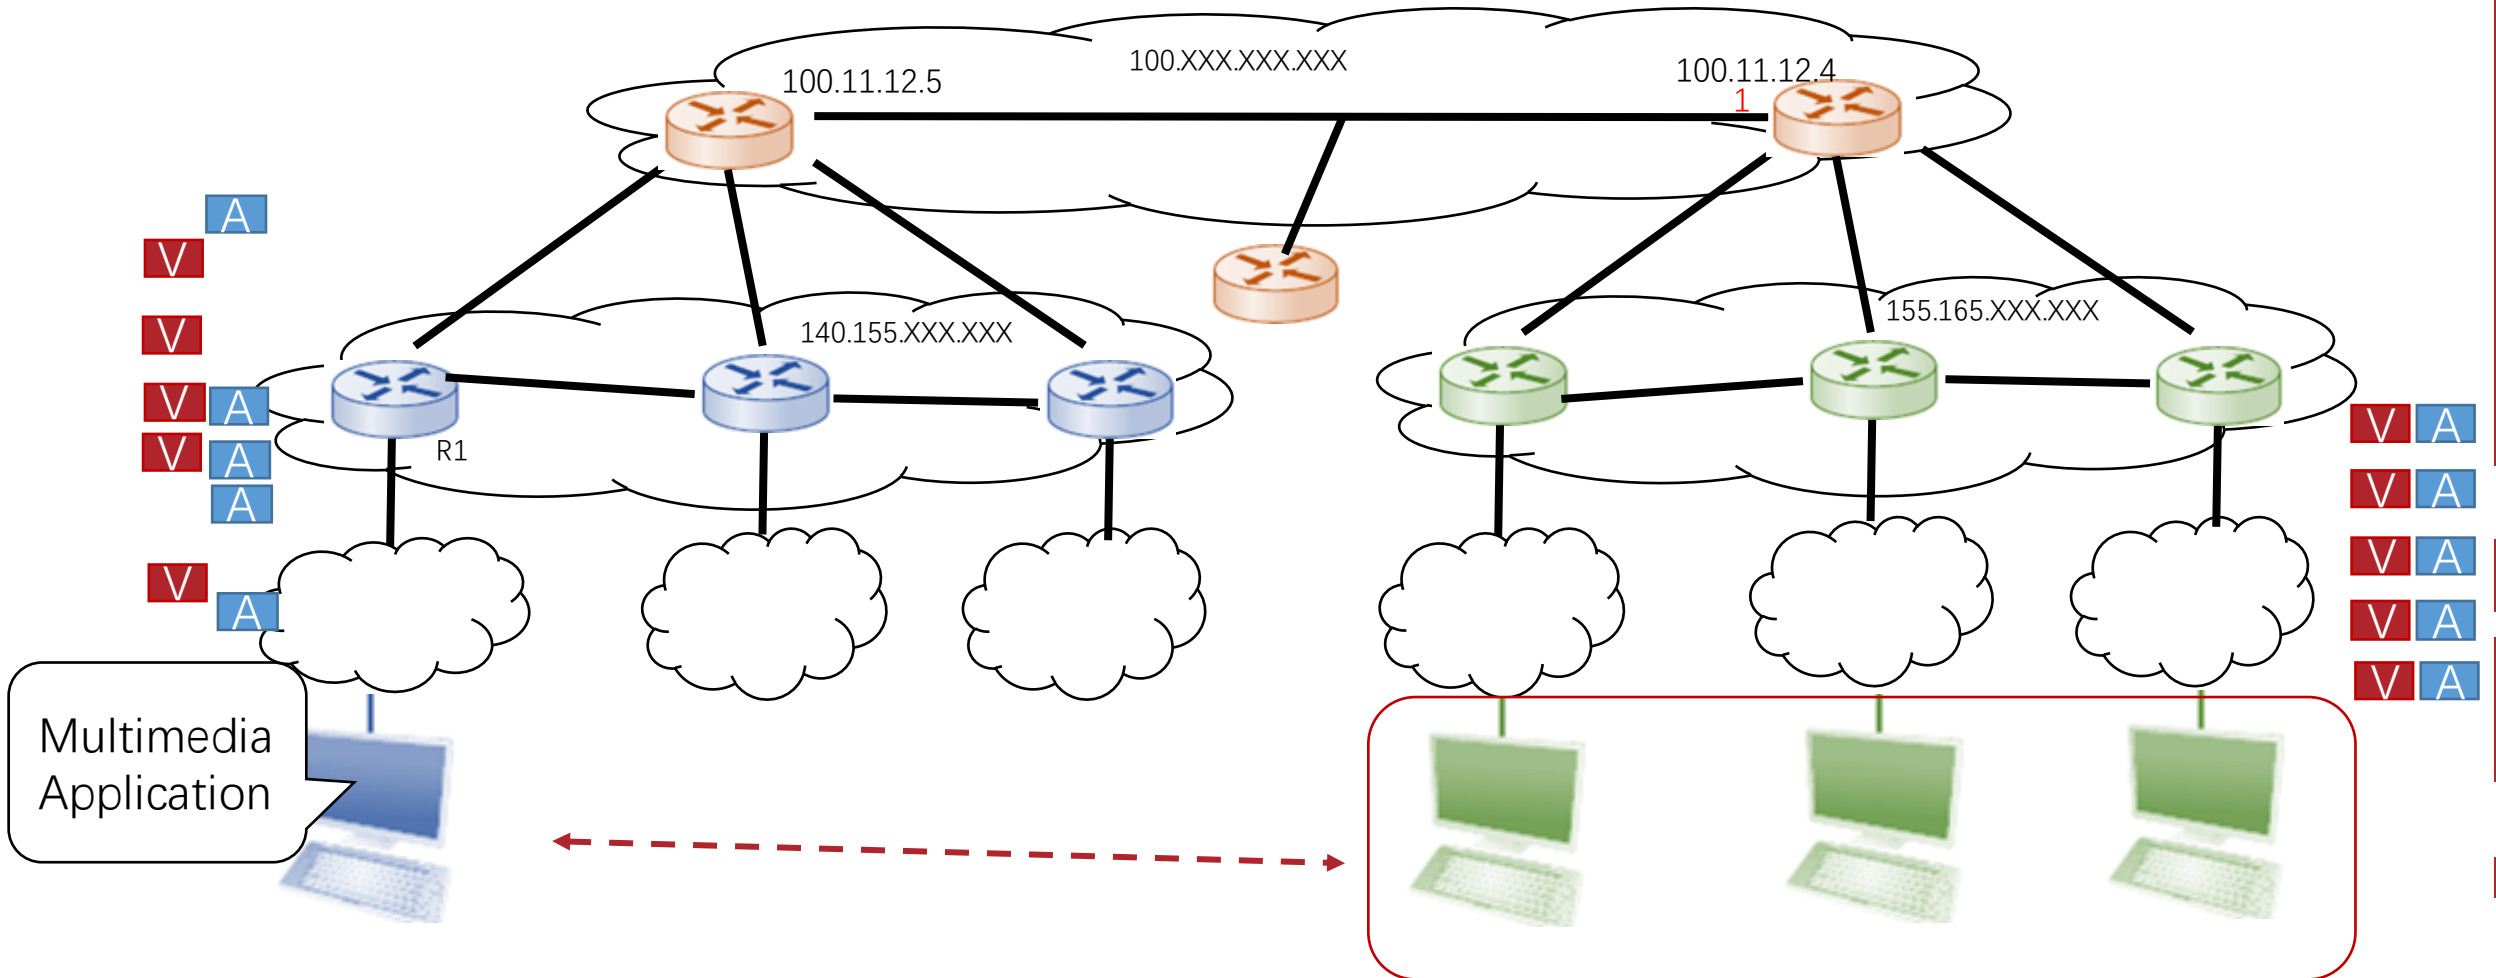
# QUIC - Parallel Streams

- Handle HOL blocking


# Outline

- TCP Fairness
- QUIC
- QoS

# Realtime Application




# Delay Profile



# Host Solution: Playback Buffer

- Buffer can be used to handle delay variance



# Quality of Service (QoS)

- Objective: to provide different service (network quality) to different applications
- Service Model
  - Best effort
  - Integrated Services (IntServ)
    - QoS supports every individual applications/flows
  - Differentiated Services (DiffServ)
    - QoS supports multiple/two classes of data or aggregated traffic



# Integrated Services (IntServ)

- Flow Specification
  - What is the flow
  - What we want to guarantee for the flow
- Admission Control
  - How network decides if it can accept the flow spec
- Resource Reservation Protocol
  - How service request gets from host to network
- Packet Classification and Scheduling
  - How routers deliver service

# Flow Specification



There are multiple options

- Specify the maximum bit rate
  - Maximum bit rate may be much higher than average
  - Reserving for the worst case is wasteful
- Specify the average bit rate
  - Network will not be able to carry bursty traffic
- **Specify the burstiness of the traffic**
  - Specify both the average rate and the burst size

# Specify Burstiness: Token Bucket

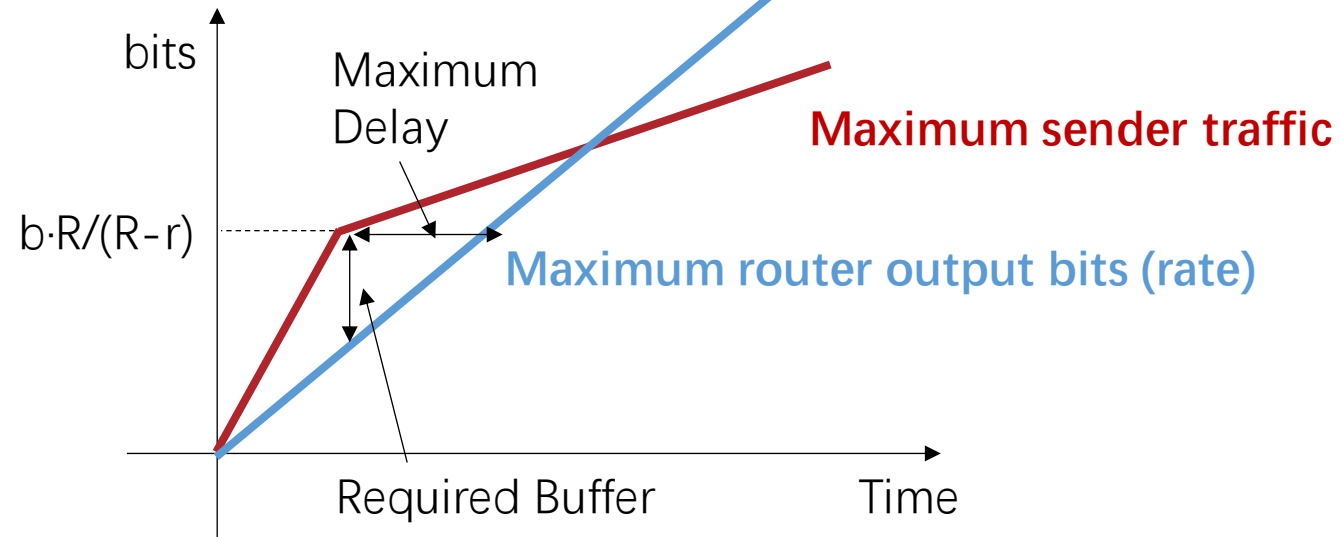
- Token Bucket: limit input to specified burst size and average rate
- Parameters:
  - $r$ : average rate, i.e., rate at which tokens fill the bucket
  - $b$ : bucket depth (limits size of burst)
  - $R$ : maximum link capacity or peak rate

Can be limited by a weighted fair queue




# Specify Burstiness: Token Bucket

- Host
  - Specify token bucket to describe its traffic
- Router
  - Allocate buffer and bandwidth to guarantee delay
  - If resource is insufficient, reject the flow




# Token Bucket


A flow can be described by multiple token bucket.




# Resource Reservation Protocol



# Resource Reservation Protocol




# Resource Reservation Protocol



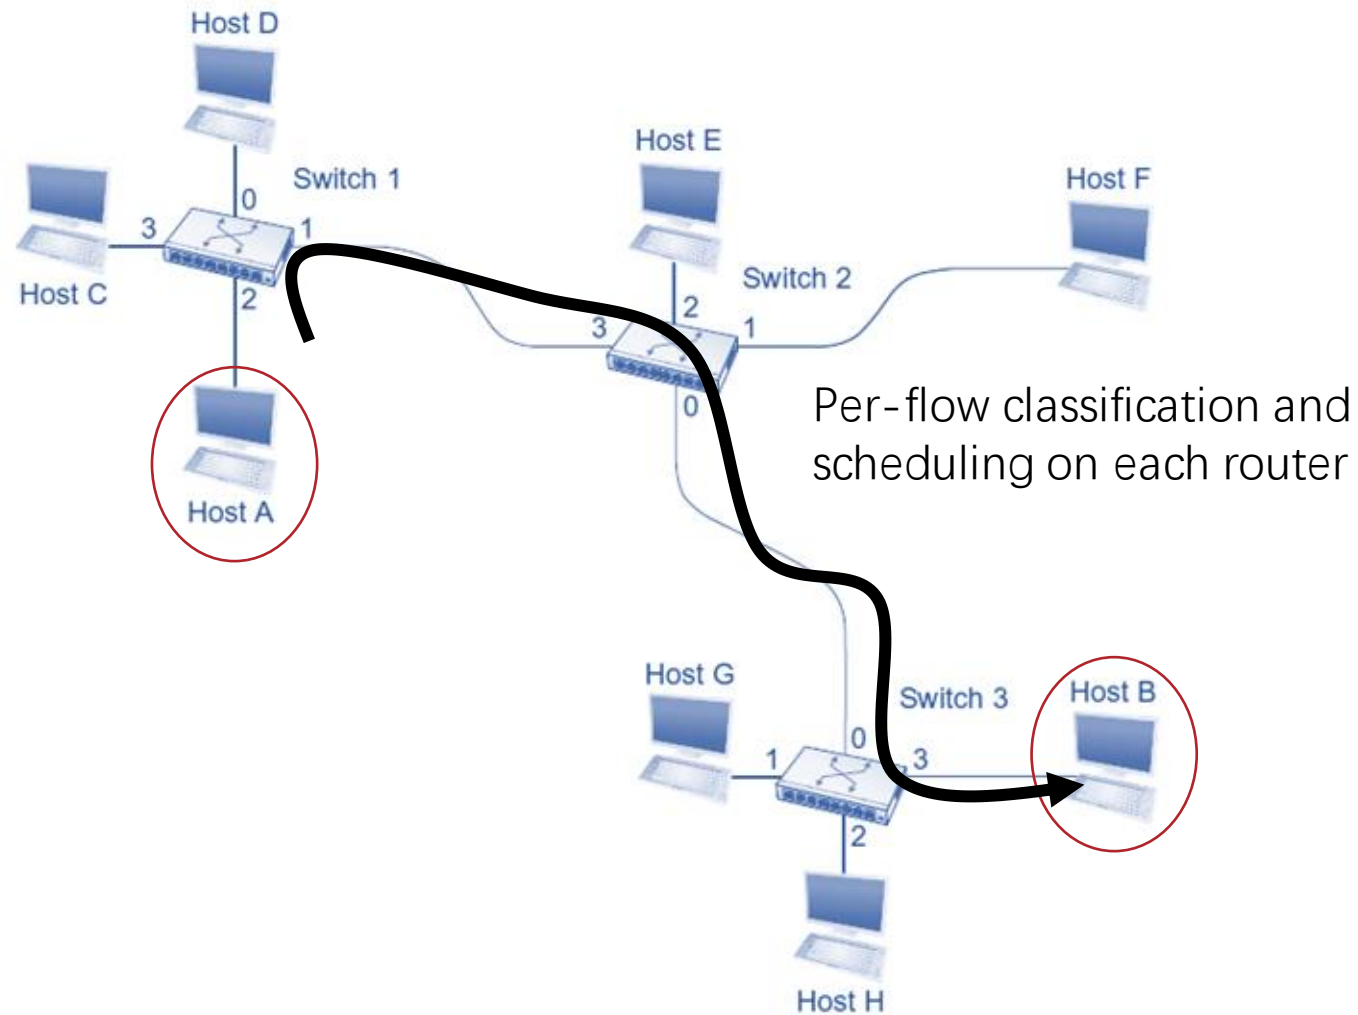
The receiver signals reservation request



# Resource Reservation Protocol



# Resource Reservation Protocol




# Packet Classification

- Classify Packets into Flows according to
  - Source Address
  - Destination Address
  - Protocol Number
  - Source Port
  - Destination Port

# Packet Scheduling

- Implementation Dependent
  - Token bucket + Fair Queue



# Scalability Issues

- Specify service for every flow is not scalable in Internet
  - Routers must keep the state of every passing flow

# Quality of Service (QoS)

- Objective: to provide different service (network quality) to different applications
- Service Model
  - Best effort
  - Integrated Services (IntServ)
    - QoS supports every individual applications/flows
  - Differentiated Services (DiffServ)
    - QoS supports multiple/two classes of data or aggregated traffic

# Differentiated Services (DiffServ)

- Problem with IntServ: scalability
  - Maintain per-flow state
  - Per-flow classification
- DiffServ Approach
  - Segregate packets into a small number of (two) classes
    - Premium
    - Other
  - Class of certain packet (state) is kept in packet header
    - ToS

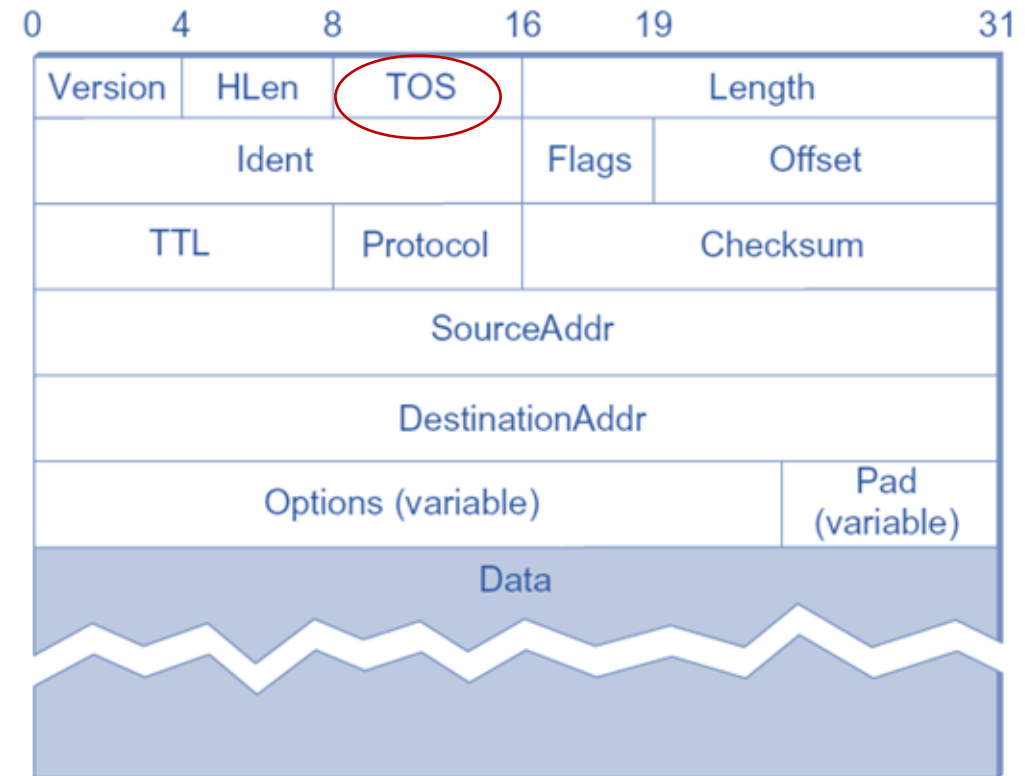
# Differentiated Services (DiffServ)

- Edge Routers
  - Set the traffic class
    - ToS field => token bucket parameters
  - Schedule traffic according to the traffic class
- Core Routers
  - Schedule traffic according to the traffic class specified by the edge router



# Per Hop Behavior

- Reuse ToS Field
  - 0-5bit: Differentiated Service Code Point (DSCP) Field
  - 6-7bit: Explicit Congestion Notification
- DSCP field encodes Per-Hop Behavior
  - Expedited Forwarding (all packets receive minimal delay & loss)
  - Assured Forwarding (packets marked with low/high drop probabilities)



# Set Packet Class

- DSCP Field in Practice
  - Edge Routers
    - Set Differentiated Service (DS) Field in IP header
      - Maybe because the user paid the ISP
  - Core Routers
    - Implement Per Hop Behavior
      - According to DS Field of packets

## Commonly used DSCP values


DSCP value	Hex value	Decimal value	Meaning	Drop probability	Equivalent IP precedence value
101 110	0x2e	46	Expedited forwarding (EF)	N/A	101 Critical
000 000	0x00	0	Best effort	N/A	000 - Routine
001 010	0x0a	10	AF11	Low	001 - Priority
001 100	0x0c	12	AF12	Medium	001 - Priority
001 110	0x0e	14	AF13	High	001 - Priority
010 010	0x12	18	AF21	Low	010 - Immediate
010 100	0x14	20	AF22	Medium	010 - Immediate
010 110	0x16	22	AF23	High	010 - Immediate
011 010	0x1a	26	AF31	Low	011 - Flash
011 100	0x1c	28	AF32	Medium	011 - Flash
011 110	0x1e	30	AF33	High	011 - Flash
100 010	0x22	34	AF41	Low	100 - Flash override
100 100	0x24	36	AF42	Medium	100 - Flash override
100 110	0x26	38	AF43	High	100 - Flash override

# Implementation of Per-Hop Behavior

- Expedited Forwarding (EF) PHB
  - Highest Priority
- Assured Forwarding (AF) PHB
  - Different levels of priorities, drop probabilities, bandwidth, etc.


# Implementation of Expedited Forwarding

- First-In-First-Out (FIFO) with Priority




# Implementation of Expedited Forwarding

- Weighted Fair Queuing (FQ)



# Implementation of Assured Forwarding


- RED with In and Out (RIO)



月基本费		<del>258元</del>	188元	
包含	国内流量	1GB		
	国内通话	800分钟		
	本地流量	本地流量无限量权益 (用满40GB后限速)		

# Network Neutrality

- Network Neutrality
  - ISPs supply non-discriminated IP connectivity





# Network Neutrality

- Opposite Counterpoint
  - ISPs only allows you to access their (often value-added) services



# Reference

- Textbook 6.5
- Some slides are adapted from [http://www-net.cs.umass.edu/kurose\\_ross/ppt.htm](http://www-net.cs.umass.edu/kurose_ross/ppt.htm) by Kurose Ross
- <https://www.chromium.org/quic>