Lecture 10: Introduction to Model Checking

Formal Verification/Validation

Grand challenge:

Automate the process as much as possible !

Analysis Techniques

- Dynamic Analysis (runtime)
  - Execute the system, possibly multiple times with different inputs
  - Check if every execution meets the desired requirement
- Static Analysis (design time)
  - Analyze the source code or the model for possible bugs
- Trade-offs
  - Dynamic analysis is incomplete, but accurate (checks real system, and bugs discovered are real bugs)
  - Static analysis can catch design bugs early !
  - Many static analysis techniques are not scalable (solution: analyze approximate versions, can lead to false warnings)

Verification Methods

- Simulation
  - Simulate the model, possibly multiple times with different inputs
  - Easy to implement, scalable, but no correctness guarantees

- Proof based
  - Construct a proof that system satisfies the invariant
  - Requires manual effort (partial automation possible)

- State-space analysis (Model checking)
  - Algorithm explores “all” reachable states to check invariants
  - Not scalable, but current tools can analyze many real-world designs (relies on many interesting theoretical advances)

Different Requirements

- Safety
  - A system always stays within “good’ states (i.e. a nothing bad ever happens)
  - Leader election: it is never the case that two nodes consider them to be leaders
  - Collision avoidance: Distance between two cars is always greater than some minimum threshold
- Liveness
  - System eventually attains its goal
  - Leader election: Each node eventually makes a decision
  - Cruise controller: Actual speed eventually equals desired speed
  - A car will always eventually reach its destination

Model Checking

Model Checking

- A domain D representing the state space of a model
- The reachable state space M for the model
- Define a subset of the state space as property φ
- Explore the whole reachable state space of a model for property violations

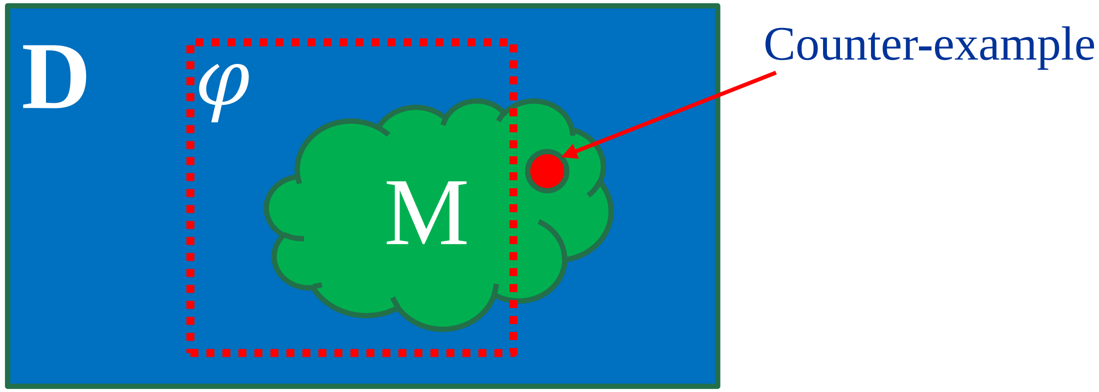

Plato vs. Diogenes

- The definition of "*human*"

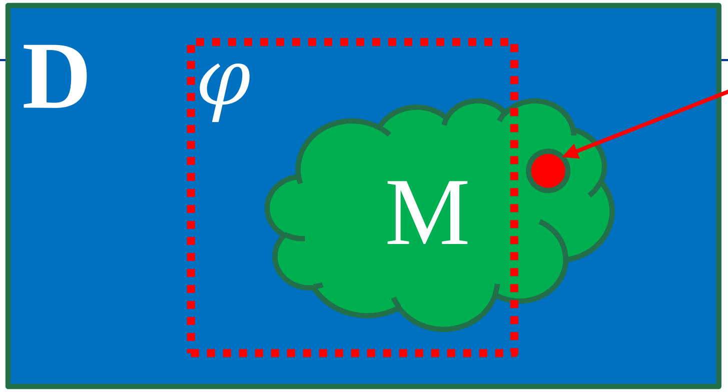

Counter-example

φ

All living creatures D

$Counter - example$

Featherless Biped

M

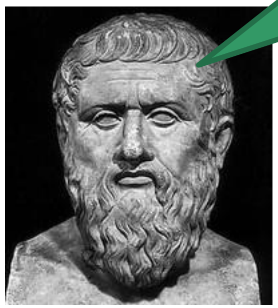

Plato

Here's Plato's human!!!!

Diogenes

Challenge

- State space explosion
  - Not every model can be model checked!!
  - i.e. Real-number (continuous) state space
  - Complex dynamics between states

- Solution
  - Simple yet expressive formalisms
  - Symbolic states/executions
  - Model abstraction/approximation

Simple yet expressive formalisms

Basic Finite State Machine (FSM)

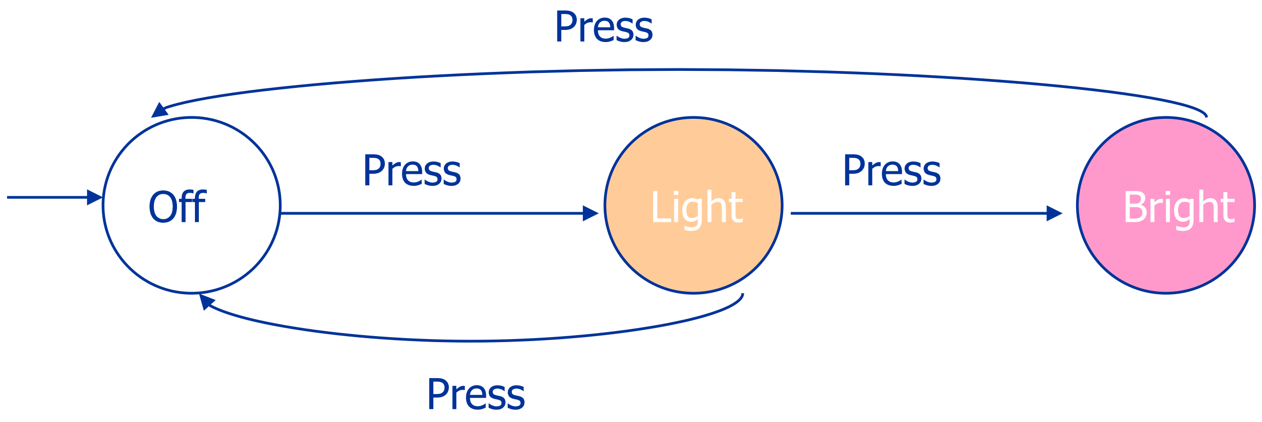

WANT: if press is issued twice quickly
then the light will get brighter; otherwise the light is turned off.

FSM with real number time

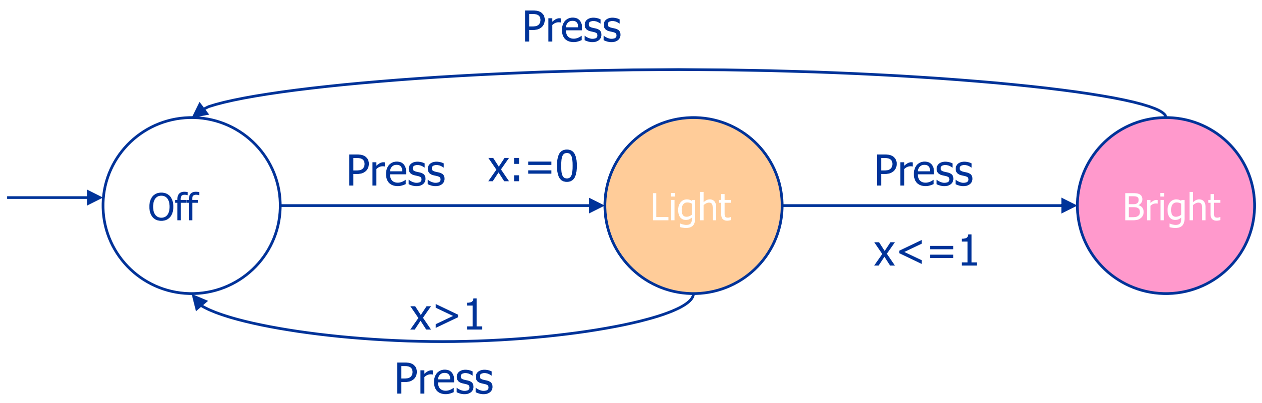

Solution: Add a real-valued clock x

Adding continuous variables to state machines

Timed Automata

Clocks: *x, y*

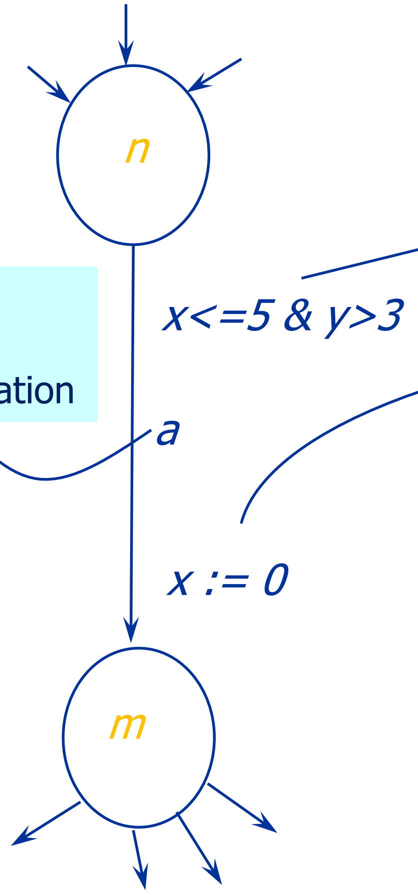

Guard
Boolean combination of comparisons with
Integer/rational bounds

Reset
Action performed on clocks

Action
used
for synchronization

State
(*location* , *x*=v , *y*=u)   where v,u are in $\mathbb{R}$

Transitions

$(n , x=2.4 , y=3.1415) \xrightarrow{a} (m , x=0 , y=3.1415)$

$(n , x=2.4 , y=3.1415) \xrightarrow{wait(1.1)} (n , x=3.5 , y=4.2415)$

Adding Invariants

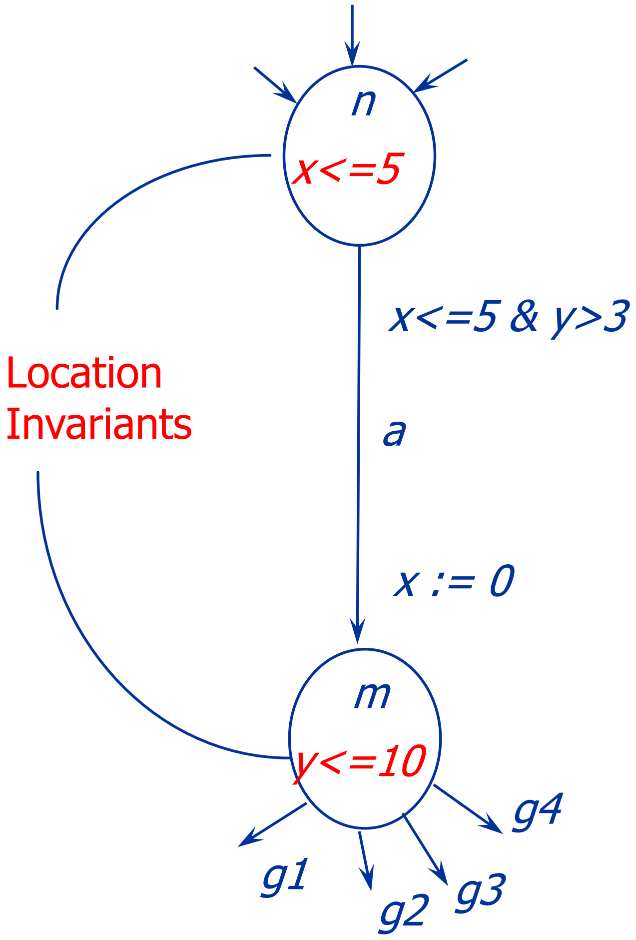

Clocks: *x, y*

Transitions

(*n* , *x*=2.4 , *y*=3.1415) $\xrightarrow{\cancel{wait(3.2)}}$

(*n* , *x*=2.4 , *y*=3.1415) $\xrightarrow{wait(1.1)}$ (*n* , *x*=3.5 , *y*=4.2415)

Invariants ensure progress!!

Timed Automata: Syntax

- A finite set *V* of locations
- A subset V^0 of initial locations
- A finite set Σ of labels (alphabet)
- A finite set *X* of clocks
- Invariant *Inv(l)* for each location: (clock constraint over *X*)
- A finite set *E* of edges. Each edge has
  - source location *l*, target location *l'*
  - label *a* in Σ (ε labels also allowed)
  - guard *g* (a clock constraint over *X*)
  - a subset λ of clocks to be reset

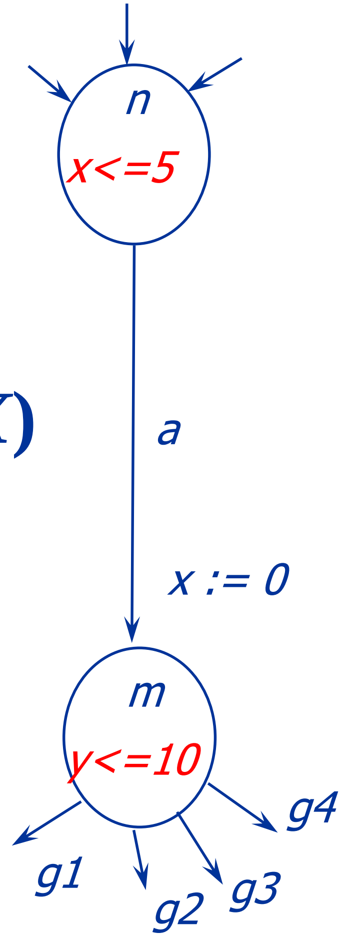

Timed Automata: Semantics

- **For a timed automaton *A*, define an infinite-state transition system *S(A)***
- **States *Q*: a state *q* is a pair *(l,v)*, where *l* is a location, and *v* is a clock vector, mapping clocks in *X* to *R*, satisfying *Inv(l)***
- ***(l,v)* is initial state if *l* is in V^0 and *v(x)=0***
- **Elapse of time transitions: for each nonnegative real number *d*, *(l,v)-d->(l,v+d)* if both *v* and *v+d* satisfy *Inv(l)***
- **Location switch transitions: *(l,v)-a->(l',v')* if there is an edge *(l,a,g,λ,l')* such that *v* satisfies *g* and *v'=v[λ:=0]***

Product Construction

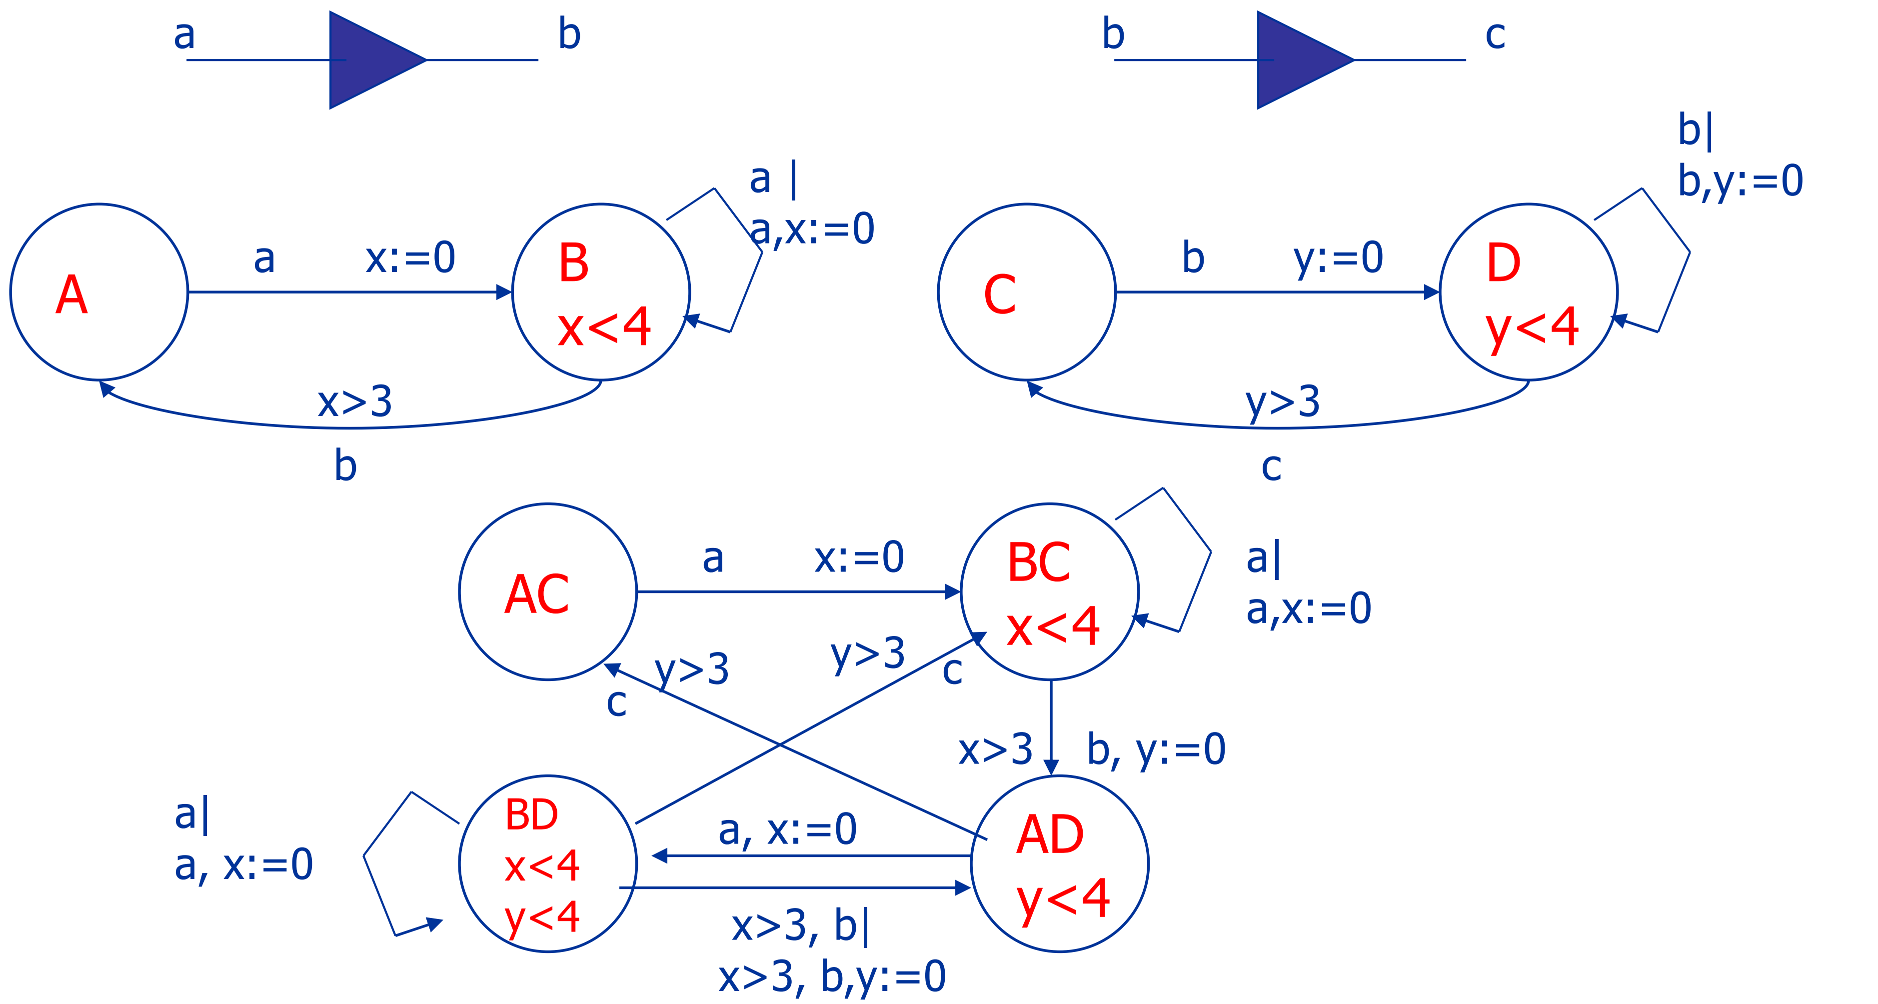

Model Checking: Forward Reachability

- Given a timed automata and a property φ
- R:=I
- Repeat
  - If R intersects $\neg\varphi$, report “yes”
  - Else if R contains Post(R), report “no”
  - Else R := R union Post(R)

Reachability for Timed Automata

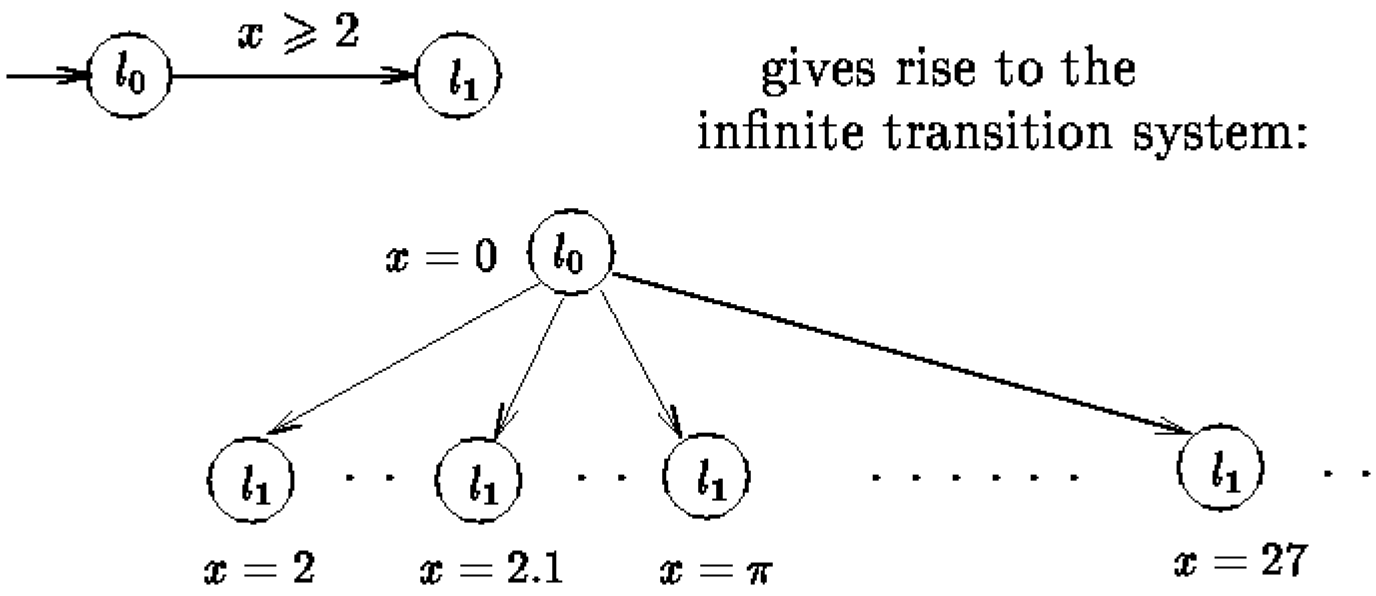

Symbolic states/executions

Finite Partitioning

Goal: To partition state-space into finitely many equivalence classes so that equivalent states exhibit similar behaviors

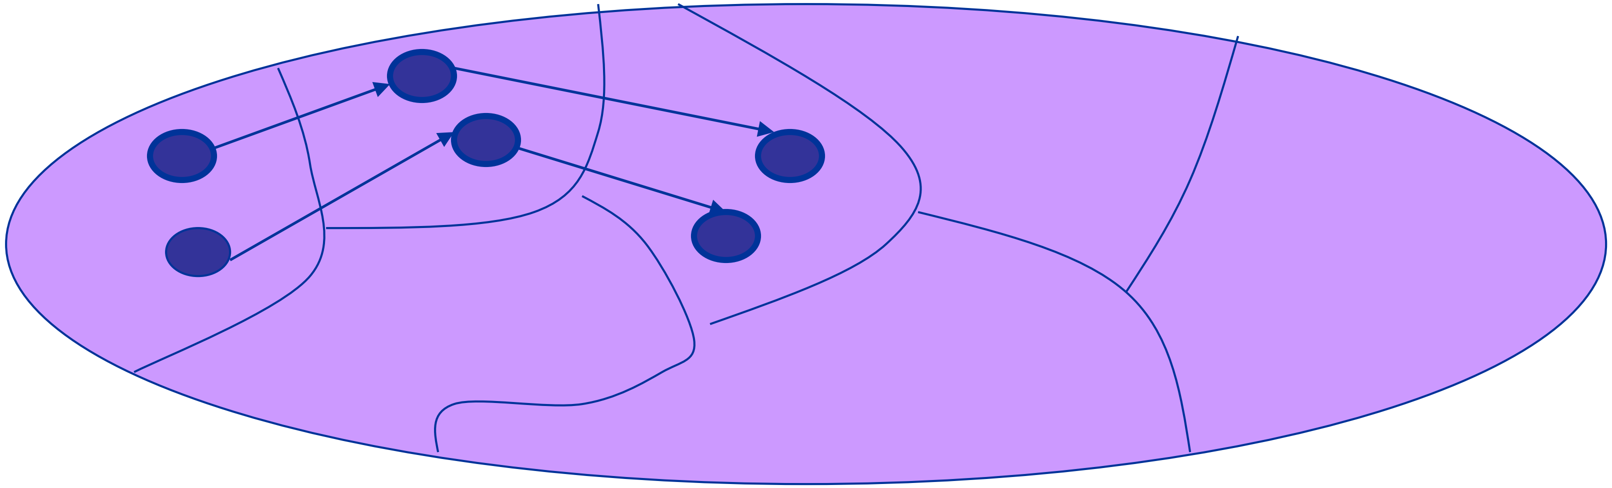

Symbolic States/executions

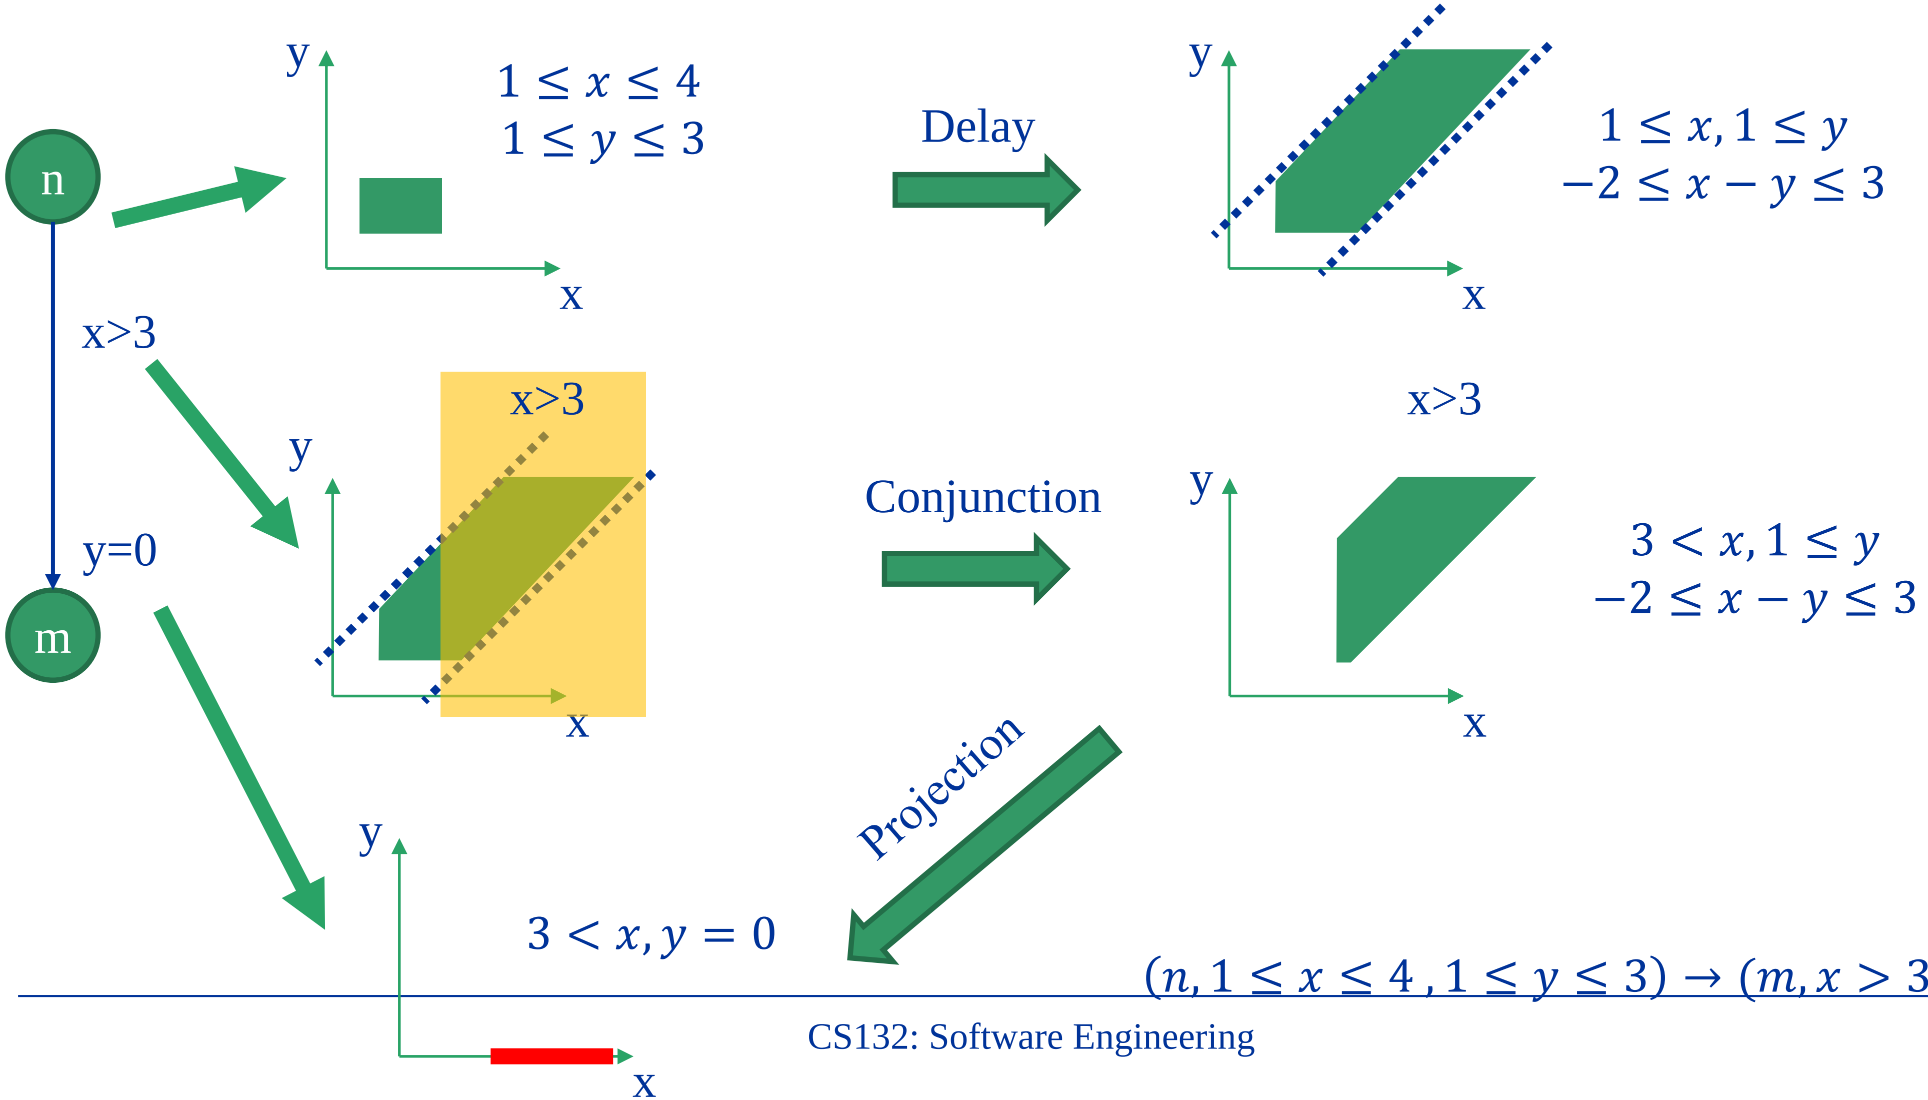

$(n, 1 \leq x \leq 4\,, 1 \leq y \leq 3) \rightarrow (m, x > 3, y = 0)$

Model abstraction/approximation

Existential Abstraction (Over-approximation)

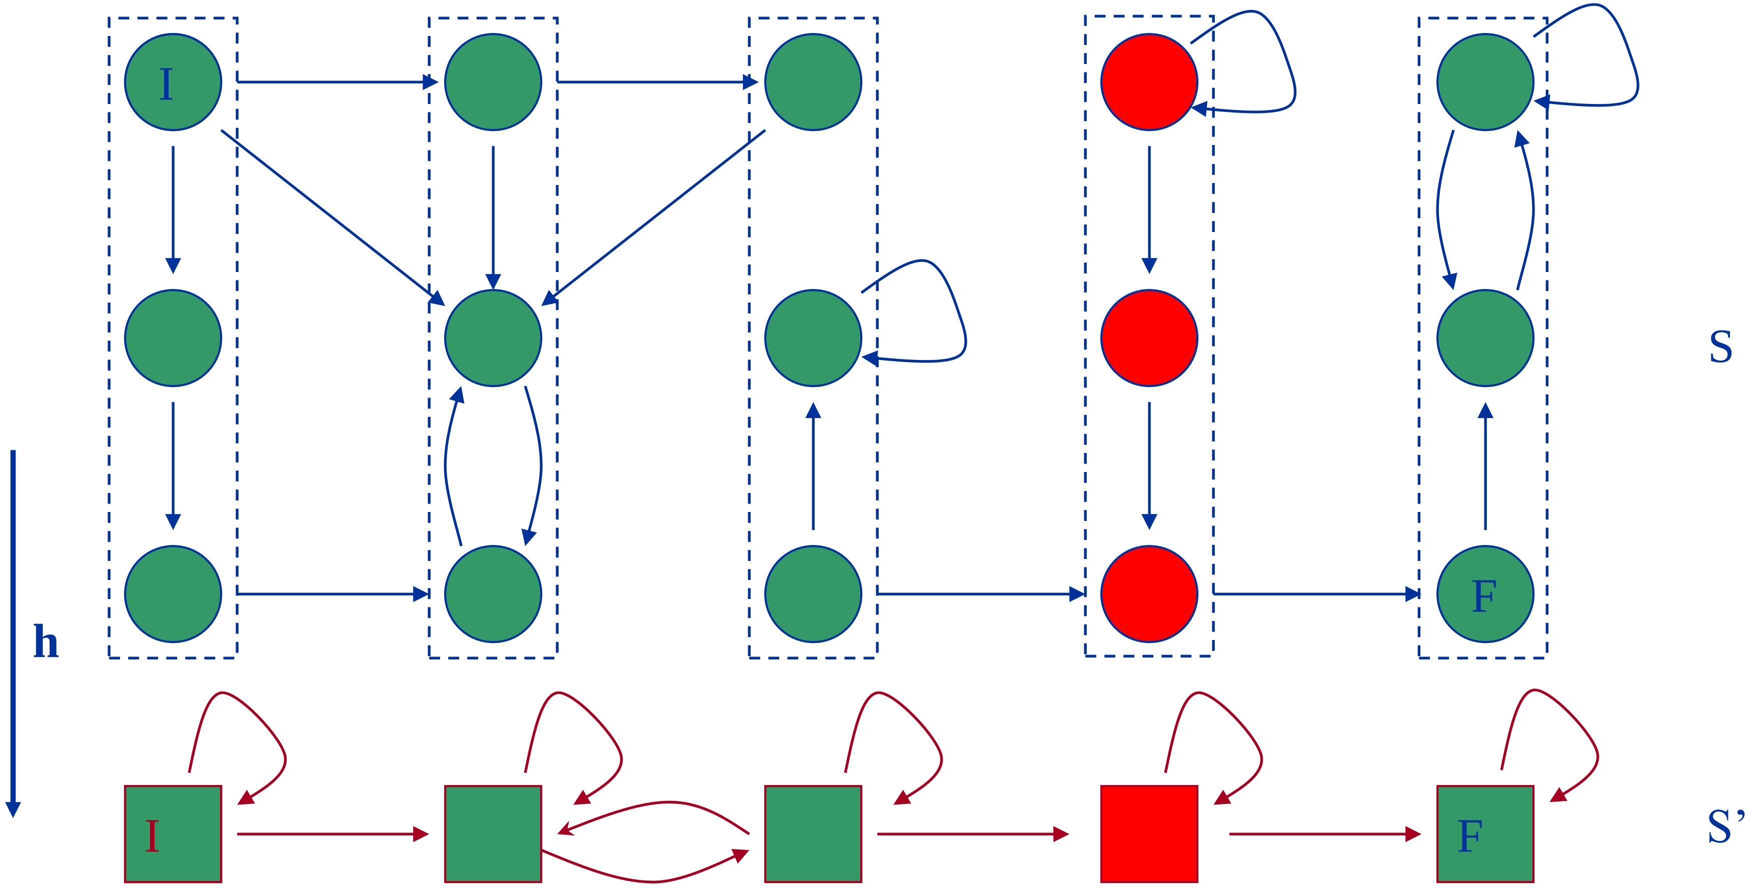

Pros and Cons of Over-approximation

- Properties satisfied by M are also satisfied by P
  - Can model check a less complex model
- M has more behaviors than P
- If a counter-example returns, it may not be a behavior of P

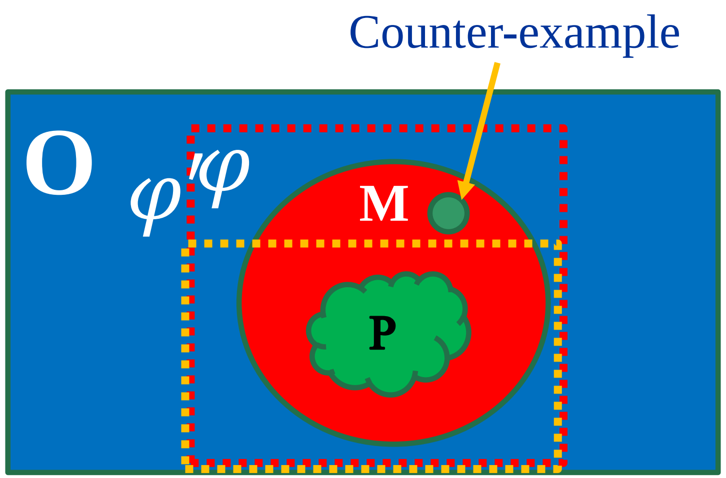

Why spurious counterexample?

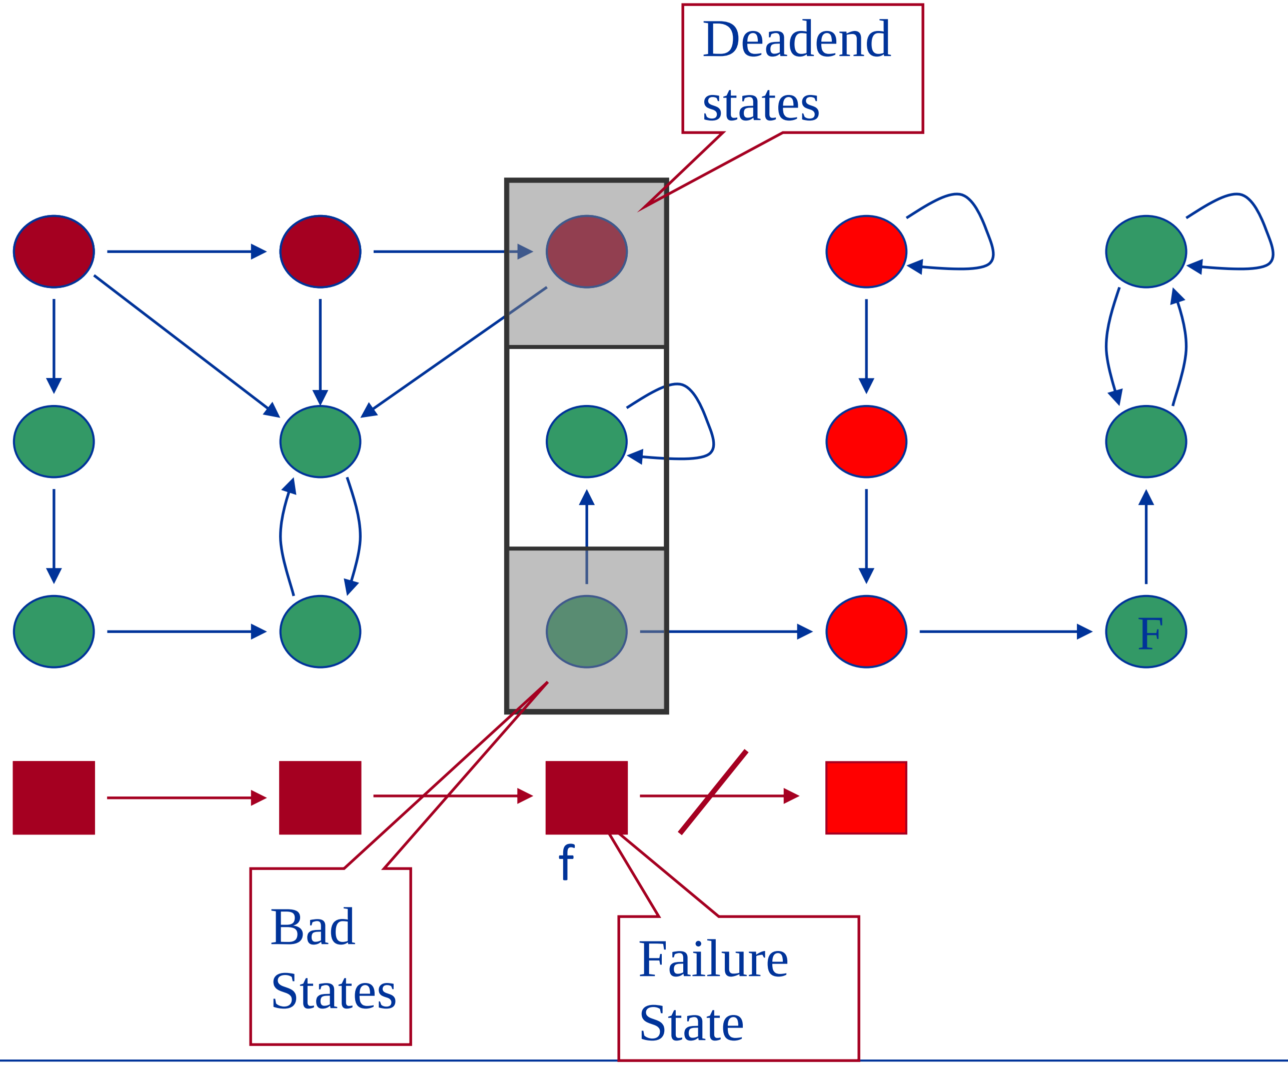

Refinement

- Problem: Deadend and Bad States are in the same abstract state.
- Solution: Refine abstraction function.
- The sets of Deadend and Bad states should be separated into different abstract states.

Refinement

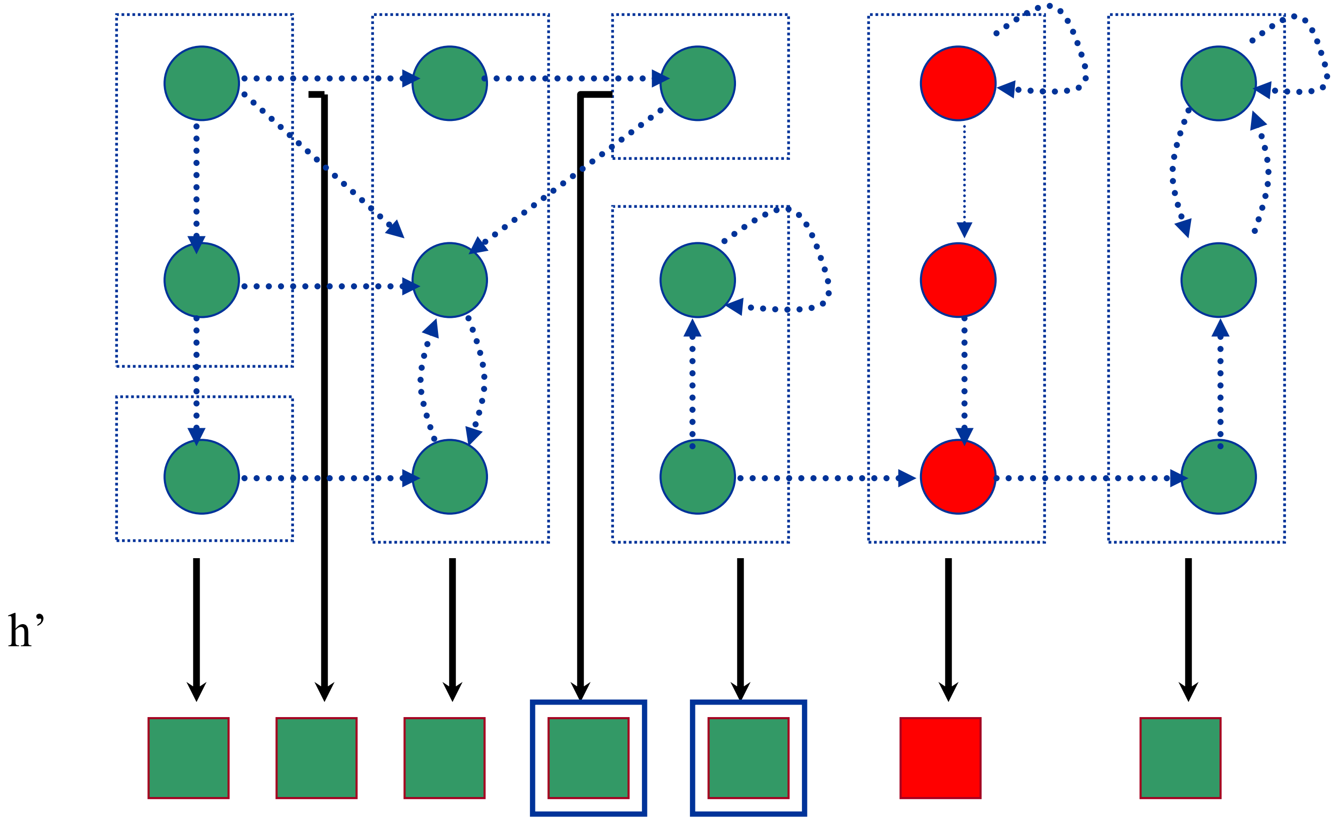

Refinement : h’

Counter-Example-Guided Abstraction and Refinement (CEGAR)

- Obtain initial abstraction

1. Model checking
2. Property satisfied -> no bugs
3. Property unsatisfied -> counter-examples
4. Check whether the CE is spurious
5. If not, bug found
6. If yes, refine the model and start from 1 again

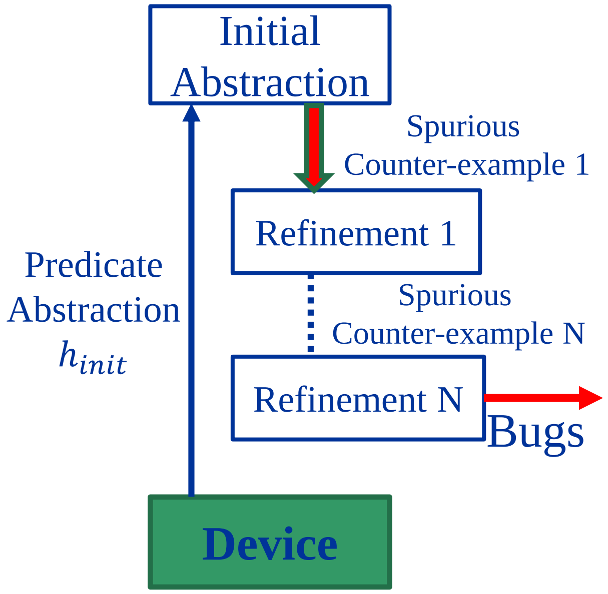

Capture Environmental Variability With Over-approximation

- Properties satisfied by M are also satisfied by P1, P2
- Behaviors not exist in P1, P2 may also be physiologically-valid
- Is this a valid counter-example?
- Need a framework to provide context for counter-examples