

PyTorch Tutorial

TA: Binbin Chen

Email: chenbb@shanghaitech.edu.cn

What is PyTorch?

- a popular machine learning framework in Python
- 2 main features:
 - N-dimensional Tensor computation on GPUs
 - Automatic differentiation for training

Tensors:

- multidimensional arrays, check with `.shape` (1D-vector, 2D-matrix)
- common operations
- device:
 - `device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')`
 - `.to(device)`
- gradient calculation:
 - `x = torch.tensor(..., requires_grad=True); z = func(x); z.backward(); x.grad`

Linear Regression:

1. Dataset Generator:

```
def synthetic_data(w, b, num_examples): #@save """ y = Xw+b+noise
""" X = torch.normal(0, 1, (num_examples, len(w))) y =
torch.matmul(X, w) + b y += torch.normal(0, 0.01, y.shape) return
X, y.reshape((-1, 1)) true_w = torch.tensor([2, -3.4]) true_b = 4.2
features, labels = synthetic_data(true_w, true_b, 1000)
```

2. Load Data

```
def data_iter(batch_size, features, labels): num_examples = len(fea
tures) indices = list(range(num_examples)) # 这些样本是随机读取的, 没有
特定的顺序 random.shuffle(indices) for i in range(0, num_examples, ba
tch_size): batch_indices = torch.tensor( indices[i: min(i + batch_s
ize, num_examples)]) yield features[batch_indices], labels[batch_in
dices]
```

```
batch_size = 10 for X, y in data_iter(batch_size, features,
labels): print(X, '\n', y) break ''' tensor([[ 0.1649, -1.1651],
[-2.0755, -1.0165], [-0.2189, 0.7607], [ 0.6833, 0.3537], [-0.2736,
-2.0485], [-0.3026, 0.9771], [ 2.4795, 0.6881], [-0.2045, -0.8509],
[-0.1353, 0.5476], [ 0.3371, -0.0479]]) tensor([[ 8.4901], [
3.5015], [ 1.1779], [ 4.3752], [10.6125], [ 0.2845], [ 6.8094], [
6.6776], [ 2.0598], [ 5.0189]]) '''
```

3. Weight Initialization

```
w = torch.normal(0, 0.01, size=(2,1), requires_grad=True) b = torch.zeros(1, requires_grad=True)
```

4. Model Definition

```
def linreg(X, w, b): #@save """线性回归模型""" return torch.matmul(X, w) + b
def squared_loss(y_hat, y): #@save """均方损失""" return (y_hat - y.reshape(y_hat.shape)) ** 2 / 2
def sgd(params, lr, batch_size): #@save """小批量随机梯度下降""" with torch.no_grad(): for param in params: param -= lr * param.grad / batch_size
param.grad.zero_()
```

5. Model Training

each epoch:

load data → model prediction → get loss → backpropagation for grads → SGD
optimize

```
lr = 0.03 num_epochs = 3 net = linreg loss = squared_loss
for epoch in range(num_epochs):
    for X, y in data_iter(batch_size, features, labels):
        l = loss(net(X, w, b), y) # X和y的小批量损失 # 因为l形状是(batch_size,1), 而不是一个标量。l中的所有元素被加到一起, # 并以此计算关于[w,b]的梯度
        l.sum().backward()
        sgd([w, b], lr, batch_size) # 使用参数的梯度更新参数
    with torch.no_grad():
        train_l = loss(net(features, w, b), labels)
    print(f'epoch {epoch + 1}, loss {float(train_l.mean()):f}')
    train_l.mean()
```

6. Test

```
print(f'w的估计误差: {true_w - w.reshape(true_w.shape)}')
print(f'b的估计误差: {true_b - b}')
```

Linear Regression - simple version:

1. Dataset Generator

2. Dataloader

```
def load_array(data_arrays, batch_size, is_train=True): #@save
    """构造一个PyTorch数据迭代器"""
    dataset = data.TensorDataset(*data_arrays)
    return data.DataLoader(dataset, batch_size, shuffle=is_train)
batch_size = 10
data_iter = load_array((features, labels), batch_size)
```

3. Model Definition

```
from torch import nn
net = nn.Sequential(nn.Linear(2, 1))
net[0].weight.data.normal_(0, 0.01)
net[0].bias.data.fill_(0)
loss = nn.MSELoss()
trainer = torch.optim.SGD(net.parameters(), lr=0.03)
```

4. Model Training

each epoch:

load data → model prediction → get loss → backpropagation for grads → SGD
optimize

```
num_epochs = 3
for epoch in range(num_epochs):
    for X, y in data_iter:
        l = loss(net(X), y)
        trainer.zero_grad()
        l.backward()
        trainer.step()
    l = loss(net(features), labels)
    print(f'epoch {epoch + 1}, loss {l:f}')
```

MLP - simple version:

0. Pre-function:

```
def train_epoch_ch3(net, train_iter, loss, updater):
    # @save # 将模型设置为训练模式
    if isinstance(net, torch.nn.Module):
        net.train()
    # 训练损失总和、训练准确度总和、样本数
    metric = Accumulator(3)
    for X, y in train_iter:
        # 计算梯度并更新参数
        y_hat = net(X)
        l = loss(y_hat, y)
        # 使用PyTorch内置的优化器和损失函数
        updater.zero_grad()
        l.mean().backward()
        updater.step()
        metric.add(float(l.sum()), accuracy(y_hat, y), y.numel())
    # 返回训练损失和训练精度
    return metric[0] / metric[2], metric[1] / metric[2]
```

```
def train_ch3(net, train_iter, test_iter, loss, num_epochs, updater):
    # @save
    for epoch in range(num_epochs):
        train_metrics = train_epoch_ch3(net, train_iter, loss, updater)
        test_acc = evaluate_accuracy(net, test_iter)
        train_loss, train_acc = train_metrics
```

1. Model Definition and Training

```
net = nn.Sequential(nn.Flatten(), nn.Linear(784, 256), nn.ReLU(), nn.Linear(256, 10))
def init_weights(m):
    if type(m) == nn.Linear:
        nn.init.normal_(m.weight, std=0.01)
net.apply(init_weights)
```

```
batch_size, lr, num_epochs = 256, 0.1, 10
loss = nn.CrossEntropyLoss(reduction='none')
optimizer = torch.optim.SGD(net.parameters(), lr=lr)
train_iter, test_iter = d2l.load_data_fashion_mnist(batch_size)
d2l.train_ch3(net, train_iter, test_iter, loss, num_epochs, optimizer)
```

Neural Networks

1. Custom Layers

```
class MLP(nn.Module):
    # 用模型参数声明层。这里，我们声明两个全连接的层
    def __init__(self):
        # 调用MLP的父类Module的构造函数来执行必要的初始化。
        # 这样，在类实例化时也可以指定其他函数参数，例如模型参数params（稍后将介绍）
        super().__init__()
        self.hidden = nn.Linear(20, 256) # 隐藏层
        self.out = nn.Linear(256, 10) # 输出层
        # 定义模型的前向传播，即如何根据输入X返回所需的模型输出
        def forward(self, X):
            # 注意，这里我们使用ReLU的函数版本，其在nn.functional模块中定义。
            return self.out(F.relu(self.hidden(X)))
```

```
class MySequential(nn.Module):
    def __init__(self, *args):
        super().__init__()
        for idx, module in enumerate(args):
            # 这里, module是Module子类的一个实例。我们把它保存在'Module'类的成员 # 变量 _modules中。_module的类型是OrderedDict
            self._modules[str(idx)] = module
        def forward(self, X):
            # OrderedDict保证了按照成员添加的顺序遍历它们
            for block in self._modules.values():
                X = block(X)
            return X
    net = MySequential(nn.Linear(20, 256), nn.ReLU(), nn.Linear(256, 10))
```

2. Parameters

```
# Access Parameters
print(net[2].state_dict())
'''
OrderedDict([('weight', tensor([[ 0.3016, -0.1901, -0.1991, -0.1220, 0.1121, -0.1424, -0.3060, 0.3400]])), ('bias', tensor([-0.0291]))])
'''
print(*[(name, param.shape) for name, param in net.named_parameters()])
'''
('0.weight', torch.Size([8, 4])) ('0.bias', torch.Size([8])) ('2.weight', torch.Size([1, 8])) ('2.bias', torch.Size([1]))
'''
```

```
# Parameter Initialization
def init_normal(m):
    if type(m) == nn.Linear:
        nn.init.normal_(m.weight, mean=0, std=0.01)
        nn.init.zeros_(m.bias)
    net.apply(init_normal)
net[0].weight.data[0], net[0].bias.data[0]
def init_xavier(m):
    if type(m) == nn.Linear:
        nn.init.xavier_uniform_(m.weight)
    def init_42(m):
        if type(m) == nn.Linear:
            nn.init.constant_(m.weight, 42)
    net[0].apply(init_xavier)
    net[2].apply(init_42)
print(net[0].weight.data[0])
print(net[2].weight.data)
```

Save & Load Models:

```
torch.save(model.state_dict(), path)
ckpt = torch.load(path)
model.load_state_dict(ckpt)
```