

Matrix Factorization

Ziping Zhao

School of Information Science and Technology
ShanghaiTech University, Shanghai, China

CS182: Introduction to Machine Learning (Spring 2023)
<http://cs182.sist.shanghaitech.edu.cn>

Outline

Introduction

Non-negative Matrix Factorization

Probabilistic Matrix Factorization

Outline


Introduction

Non-negative Matrix Factorization

Probabilistic Matrix Factorization

Introduction

- ▶ **Matrix factorization (MF)** (or **matrix decomposition**) refers to techniques that approximate a matrix by the product of two or more (smaller) matrices.
- ▶ PCA as a dimensionality reduction technique may also be seen as a matrix factorization method that is subject to an orthogonality constraint.
- ▶ MF is a generalization of many methods (e.g., PCA, SVD, QR, CUR, Truncated SVD, etc.)
- ▶ Depending on the applications, other constraints may also be enforced to give different matrix factorization methods.
- ▶ Two methods considered here:
 - **Non-negative matrix factorization (NMF)**
 - **Probabilistic matrix factorization (PMF)**



Outline

Introduction

Non-negative Matrix Factorization

Probabilistic Matrix Factorization

Non-negative Matrix Factorization

Non-negative Matrix Factorization

- ▶ Non-negative matrix factorization (NMF) was first proposed as a method for learning **parts-based representations** in visual perception tasks, as opposed to other methods such as vector quantization (VQ) and PCA which learn holistic, not parts-based representations.
- ▶ NMF algorithms are very useful in a wide variety of machine learning applications as a dimensionality reduction or feature extraction method.
- ▶ It enforces **non-negativity constraints** on the factor matrices, i.e., all entries in the factor matrices are non-negative.
- ▶ The non-negativity constraints make the method biologically more plausible, e.g., the firing rates of biological neurons and the synaptic strengths between neurons are non-negative.

Matrix Factorization for Representing Facial Images

- ▶ A database consists of m facial images.
- ▶ Each image consists of n nonnegative pixels.
- ▶ The whole database can be represented by a matrix \mathbf{V} of size $n \times m$.
- ▶ \mathbf{V} is approximated by two factor matrices \mathbf{W} (basis elements) and \mathbf{H} (encodings) of sizes $n \times r$ and $r \times m$, respectively:

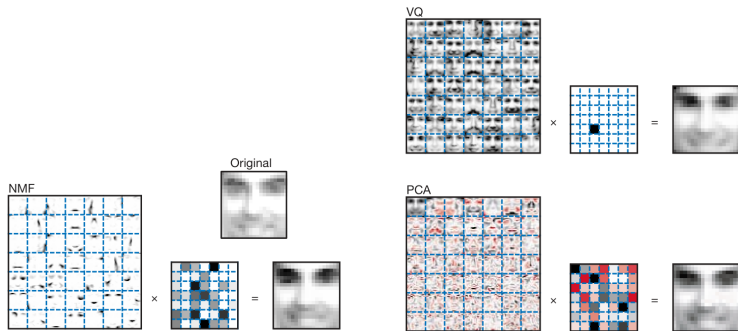
$$\mathbf{V} \approx \mathbf{WH}$$

where the **rank** r of the factorization is generally chosen such that

$$(n + m)r < nm$$

to give a **compressed form** of the data in \mathbf{V} .


NMF vs. VQ and PCA I



- ▶ Each method has learned r ($= 49 = 7 \times 7$) basis images.
- ▶ Each facial image is approximately represented by a linear combination of the r basis images.
- ▶ Positive values are shown in black and negative values in red.

NMF vs. VQ and PCA II

- ▶ Each of the basis images learned by VQ is a whole face.
- ▶ Each of the basis images learned by PCA is called an **eigenface**. Some eigenfaces resemble distorted versions of whole faces.
- ▶ Unlike the holistic basis images learned by VQ and PCA, those learned by NMF are **localized features** that correspond to parts of faces, i.e., parts-based representations, such as the nose, mouth, moustaches, and eyes of a face.



NMF vs. VQ and PCA III

- ▶ The combinations in NMF can only be additive due to the non-negative constraints on the matrices.
- ▶ This offers an intuitive explanation as to why NMF learns a parts-based representation (in contrast to the holistic representations of faces resulting from PCA and VQ), as parts of the face are additively combined to create a whole.
- ▶ It is also important to note that the NMF basis images and encodings contain several vanishing coefficients, meaning that the basis images and encodings are **sparse**, which is crucial for parts-based representation.

NMF Problem Formulation

- ▶ Given a non-negative matrix \mathbf{V} , find non-negative matrices \mathbf{W} and \mathbf{H} such that

$$\mathbf{V} \approx \mathbf{WH}$$

- ▶ Each column vector \mathbf{v} of \mathbf{V} is approximated by the corresponding column vector \mathbf{h} of \mathbf{H} :

$$\mathbf{v} \approx \mathbf{Wh}$$

I.e., each vector \mathbf{v} , corresponding to one facial image, is approximated by a **linear combination** of the columns of \mathbf{W} where the components of \mathbf{h} are the **weights** of the linear combination.

Cost Functions for Optimization Problem

- ▶ The optimization problem is to minimize the difference between \mathbf{V} and \mathbf{WH} which is quantified by some **cost function**.
- ▶ Two choices:

- **Euclidean distance** between \mathbf{A} and \mathbf{B} (as Frobenius norm of $\mathbf{A} - \mathbf{B}$):

$$\|\mathbf{A} - \mathbf{B}\|_F^2 = \sum_{ij} (A_{ij} - B_{ij})^2$$

- Variant of **Kullback-Leibler (KL) divergence** of \mathbf{A} with respect to \mathbf{B} :

$$D(\mathbf{A} \parallel \mathbf{B}) = \sum_{ij} \left(A_{ij} \log \frac{A_{ij}}{B_{ij}} - A_{ij} + B_{ij} \right)$$

which reduces to the KL divergence when $\sum_{ij} A_{ij} = \sum_{ij} B_{ij} = 1$. This objective can be derived by treating NMF like a probabilistic generative model.

- ▶ Both cost functions have a lower bound of zero and are equal to zero if and only if $\mathbf{A} = \mathbf{B}$.

Optimization Problem I

► Optimization problem:

$$\begin{array}{ll} \underset{\mathbf{W}, \mathbf{H}}{\text{minimize}} & \|\mathbf{V} - \mathbf{WH}\|_F^2 \\ \text{subject to} & \mathbf{W}, \mathbf{H} \geq \mathbf{0} \end{array} \quad \text{or} \quad \begin{array}{ll} \underset{\mathbf{W}, \mathbf{H}}{\text{minimize}} & D(\mathbf{V} \parallel \mathbf{WH}) \\ \text{subject to} & \mathbf{W}, \mathbf{H} \geq \mathbf{0} \end{array}$$

- As opposed to the unconstrained problem which can be solved efficiently using the SVD, NMF is NP-hard in general.
- Issues:
 - It is not guaranteed to find a single unique decomposition (in general, there might be many schemes for defining sets of basis elements). It can be tackled using other priors on \mathbf{W} and \mathbf{H} and adding proper regularization terms in the objective function.
 - It is hard to know how to choose the factorization rank r .
- The objective functions are **convex** in \mathbf{W} only or \mathbf{H} only, but **not convex** in both variables together.

Optimization Problem II

- ▶ Gradient methods such as gradient descent and conjugate gradient may be applied to find **local minima** for the optimization problem, but their convergence is very sensitive to the choice of step size.
- ▶ Almost all NMF algorithms use a **two-block coordinate descent** scheme (exact or inexact), that is, they optimize alternatively over one of the two factors, **W** or **H**, while keeping the other fixed, since the subproblem in one factor is convex.
- ▶ More precisely, it is a **nonnegative least squares (NNLS)** problem, a type of constrained least squares problems.
- ▶ Many algorithms exist to solve the NNLS problem; and NMF algorithms based on two-block coordinate descent differ by which NNLS algorithm is used.
- ▶ Some NNLS algorithms that can be used include multiplicative updates, alternating nonnegative least squares, hierarchical alternating least squares, etc.

Multiplicative Update Rules

- ▶ The Euclidean distance $\|\mathbf{V} - \mathbf{WH}\|_F^2$ is **nonincreasing** under the following **multiplicative update rules**:

$$H_{a\mu} \leftarrow H_{a\mu} \frac{(\mathbf{W}^T \mathbf{V})_{a\mu}}{(\mathbf{W}^T \mathbf{WH})_{a\mu}} \quad W_{ia} \leftarrow W_{ia} \frac{(\mathbf{VH}^T)_{ia}}{(\mathbf{WHH}^T)_{ia}}$$

The Euclidean distance is invariant under these updates if and only if \mathbf{W} and \mathbf{H} are at a stationary point of the distance.

- ▶ The divergence $D(\mathbf{V} \parallel \mathbf{WH})$ is **nonincreasing** under the following **multiplicative update rules**:

$$H_{a\mu} \leftarrow H_{a\mu} \frac{\sum_i W_{ia} V_{i\mu} / (\mathbf{WH})_{i\mu}}{\sum_k W_{ka}} \quad W_{ia} \leftarrow W_{ia} \frac{\sum_\mu H_{a\mu} V_{i\mu} / (\mathbf{WH})_{i\mu}}{\sum_v H_{av}}$$

The divergence is invariant under these updates if and only if \mathbf{W} and \mathbf{H} are at a stationary point of the divergence.

Convergence Guarantee

- ▶ Each update involves a multiplicative factor which is equal to one (i.e., no update) when $\mathbf{V} = \mathbf{WH}$. Thus perfect reconstruction is necessarily a **fixed point** of the update rules.
- ▶ **Convergence proofs** for the two cost functions exist and are similar to that for the EM algorithm by making use of an **auxiliary function**.
- ▶ Using the multiplicative update rules, **locally optimal solutions** to the optimization problem are guaranteed.

Algorithm Initializations

- ▶ You can initialize \mathbf{W} and \mathbf{H} randomly, but there are also alternate strategies designed to give better initial estimates in the hope of converging more rapidly to a good solution:
 - Using some clustering method, and make the cluster means of the top r clusters as the columns of \mathbf{W} , and \mathbf{H} as a scaling of the cluster indicator matrix (which elements belong to which cluster).
 - Finding the best rank- r approximation of \mathbf{V} using SVD and using this to initialize \mathbf{W} and \mathbf{H} .
 - Picking r columns of \mathbf{V} and just using those as the initial values for \mathbf{W} .

More Applications Based on NMF

- ▶ The applications of NMF for feature extraction are not just limited to facial images in [image processing](#).
- ▶ It can be used in [text mining](#) for semantic analysis on text from the Encyclopedia, which is called latent semantic indexing or latent semantic analysis.
 - \mathbf{V} is a sample of m documents each using a bag of words representation with n words, each factor may be one topic or concept written using a certain subset of words and each document is a certain combination of such factors.
 - NMF identifies topics and simultaneously classifies the documents among these different topics.
- ▶ It can also be used for hyperspectral imaging, calcium imaging (of neurons), biological data mining, etc.

Outline

Introduction

Non-negative Matrix Factorization

Probabilistic Matrix Factorization

Recommendation Systems

- ▶ Recommendation systems (a.k.a. recommender systems) predict the preferences of users for items (e.g., books, movies) and make recommendations accordingly.
- ▶ Examples of recommendation systems:
 - Ordering news articles to online newspaper readers based on prediction of reader interests.
 - Ordering suggestions to customers of an online retailer about what they might like to buy based on their past history of purchases and/or product searches.
- ▶ Two main approaches:
 - Content-based systems recommend items based on properties of the items.
 - Collaborative filtering (CF) systems recommend items based on similarity measures between users and/or items.

Collaborative Filtering

- ▶ The goal of CF is to infer user preferences for items given a large but incomplete collection of preferences for many users.
- ▶ An illustrative example:
 - Suppose we infer from data that most users who like 'Star Wars' also like 'Lord of the Rings' and dislike 'Dune'.
 - Then, if a user watched and liked 'Star Wars', we would recommend him/her 'Lord of the Rings' but not 'Dune'.
- ▶ Preferences can be explicit or implicit:
 - **Explicit preferences:** e.g., ratings given to items by users
 - **Implicit preferences:** e.g., which items were rented or bought by users.

Content-based vs. CF Approaches

- ▶ The **content-based approach** makes recommendations based on item content.
 - E.g., for a movie: genre, actors, director, length, language, etc.
 - Can be used to recommend new items for which no ratings are available yet.
 - Does not perform as well as CF in most cases.
- ▶ The **CF approach** does not look at item content.
 - Preferences are inferred from rating patterns alone.
 - Cannot recommend new items – they all look the same to the system.
 - Very effective when a sufficient amount of data is available.
- ▶ It is possible to combine the two approaches in different ways.


CF as Matrix Completion

- ▶ CF can be viewed as a **matrix completion** problem.

	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6
Alice	***				*****	
Bob	**	****			*	*****
Charles				****		****
David		****	****		*****	
Eva	*****			***		**
Fred	*	*****	****		****	

- ▶ Problem: given a user-item **rating matrix** with only a small subset of entries present, fill in (some of) the missing entries.
- ▶ An effective way to solve the matrix completion problem is by **matrix factorization**, which approximates the rating matrix by the product of two smaller matrices.

Matrix Factorization: Notation



- ▶ Suppose we have N users and M movies.
- ▶ Let \mathbf{R} be a **rating matrix** of integer rating values from 1 to K with R_{ij} denoting the rating of user i for movie j , and $\mathbf{U} \in \mathbb{R}^{D \times N}$ and $\mathbf{V} \in \mathbb{R}^{D \times M}$ be the **latent user feature matrix** and **latent movie feature matrix**, respectively.

A Non-probabilistic View

- ▶ To predict the rating given by user i to movie j , we simply compute the dot product between the corresponding latent feature vectors \mathbf{U}_i and \mathbf{V}_j for user i and movie j , respectively:


$$\hat{R}_{ij} = \mathbf{U}_i^T \mathbf{V}_j = \sum_k U_{ik} V_{jk}$$

- ▶ Intuition: for each user, we predict a movie rating by giving the movie feature vector to a linear model.
 - The movie feature vector can be viewed as the **input**.
 - The user feature vector can be viewed as the **weight vector**.
 - The predicted rating is the **output**.
 - Unlike in linear regression where the inputs are fixed and the weights are learned, we learn **both** the inputs and the weights here (by minimizing the squared error).
 - Note that the model is **symmetric** in the users and movies.

PMF: Generative Model I

- ▶ Probabilistic matrix factorization (PMF) is a simple probabilistic **linear model** with **Gaussian observation noise**.
- ▶ Given the feature vectors for user i and movie j , the distribution of the corresponding rating is:

$$p(R_{ij} \mid \mathbf{U}_i, \mathbf{V}_j, \sigma^2) = \mathcal{N}(R_{ij} \mid \mathbf{U}_i^T \mathbf{V}_j, \sigma^2)$$




PMF: Generative Model II

- Distribution over all observed ratings:

$$p(\mathbf{R} \mid \mathbf{U}, \mathbf{V}, \sigma^2) = \prod_{i=1}^N \prod_{j=1}^M \left[\mathcal{N}(R_{ij} \mid \mathbf{U}_i^T \mathbf{V}_j, \sigma^2) \right]^{\mathbf{I}_{ij}}$$

where $\mathbf{I}_{ij} = 1$ if user i has rated movie j and is 0 otherwise.



- The user and movie feature vectors are given **zero-mean spherical Gaussian priors**:

$$p(\mathbf{U} \mid \sigma_U^2) = \prod_{i=1}^N \mathcal{N}(\mathbf{U}_i \mid \mathbf{0}, \sigma_U^2 \mathbf{I}), \quad p(\mathbf{V} \mid \sigma_V^2) = \prod_{j=1}^M \mathcal{N}(\mathbf{V}_j \mid \mathbf{0}, \sigma_V^2 \mathbf{I})$$

PMF: Learning I

- ▶ **MAP estimation**: maximizing the log-posterior over the user and movie features with fixed hyperparameters.
- ▶ Equivalent optimization problem (minimizing the sum of squared errors with **quadratic regularization terms**):

$$\underset{\mathbf{U}, \mathbf{V}}{\text{minimize}} \quad E = \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^M \mathbf{I}_{ij} (R_{ij} - \mathbf{u}_i^T \mathbf{v}_j)^2 + \frac{\lambda_U}{2} \sum_{i=1}^N \|\mathbf{u}_i\|_2^2 + \frac{\lambda_V}{2} \sum_{j=1}^M \|\mathbf{v}_j\|_2^2$$

where $\lambda_U = \sigma^2 / \sigma_U^2$ and $\lambda_V = \sigma^2 / \sigma_V^2$.

PMF: Learning II

- ▶ The error function E is **convex** in \mathbf{U} only or \mathbf{V} only, but **not convex** in both \mathbf{U} and \mathbf{V} together.
- ▶ **Block coordinate descent** alternates between two **convex** optimization subproblems:
 - Minimizing E over \mathbf{U} with \mathbf{V} fixed
 - Minimizing E over \mathbf{V} with \mathbf{U} fixed.
- ▶ The algorithm takes time **linear** in the number of observed ratings.

Slight Variation

- ▶ Using a simple linear-Gaussian model as above can lead to predicted ratings outside the valid range.
- ▶ To overcome this problem, the dot product of the user and movie feature vectors is passed through the **logistic function**

$$g(x) = \frac{1}{1 + \exp(-x)}$$


- ▶ Distribution over all observed ratings:

$$p(\mathbf{R} \mid \mathbf{U}, \mathbf{V}, \sigma^2) = \prod_{i=1}^N \prod_{j=1}^M \left[\mathcal{N}(R_{ij} \mid g(\mathbf{U}_i^T \mathbf{V}_j), \sigma^2) \right]^{I_{ij}}$$

- ▶ The ratings $1, \dots, K$ are linearly mapped to the interval $[0, 1]$ so that the range of valid rating values matches the range of predictions made by the model.

Extension with Automatic Complexity Control I

- ▶ **Model complexity** is controlled by the noise variance σ^2 and the parameters of the priors σ_U^2 and σ_V^2 in standard PMF.
- ▶ Approach: find MAP estimates for the **hyperparameters** Θ_U and Θ_V after introducing priors for them.



Extension with Automatic Complexity Control II

- ▶ Learning: find MAP estimates of the parameters and hyperparameters by maximizing the log-posterior:

$$\begin{aligned} & \log p(\mathbf{U}, \mathbf{V}, \sigma^2, \boldsymbol{\Theta}_U, \boldsymbol{\Theta}_V \mid \mathbf{R}) \\ &= \log p(\mathbf{R} \mid \mathbf{U}, \mathbf{V}, \sigma^2) + \log p(\mathbf{U} \mid \boldsymbol{\Theta}_U) + \log p(\mathbf{V} \mid \boldsymbol{\Theta}_V) \\ & \quad + \log p(\boldsymbol{\Theta}_U) + \log p(\boldsymbol{\Theta}_V) + \text{const.} \end{aligned}$$

- ▶ Essentially this is like choosing the hyperparameters λ_U and λ_V in the standard PMF automatically.
- ▶ Automatic selection of the hyperparameter values works considerably better than the manual approach using a validation set.