



CS240 Algorithm Design and Analysis

Lecture 20

Randomized Algorithms (Cont.)


Quan Li
Fall 2023
2023.12.19



Max Cut



- We studied the Min-Cut problem, which is closely related to finding the max flow in a network.
- Max-Cut is the opposite of Min-Cut.
- Given a graph G , split vertices into two sides to maximize the number of edges between the sides.





Max Cut



- Unlike Min-Cut, Max-Cut is NP-complete.
- We'll give a very simple randomized Monte Carlo 2-approximation algorithm.
 - Monte Carlo means the algorithm sometimes returns the wrong answer, i.e., a cut that's not a 2-approximation.
 - Monte Carlo also means the algorithm always runs in a fixed amount of time.
 - ❖ Put each node in a random side with probability $\frac{1}{2}$.






Correctness



- **Lemma** In a graph with e edges, the algorithm produces a cut with expected size $e/2$.
- **Proof** Let X be a random variable equal to the size of the cut. We want to bound $E[X]$.
 - For each edge e , let X_e be the indicator variable of whether e is in the cut.
 - i.e., $X_e=1$ if e is in the cut and 0 otherwise.
 - $X = \sum_e X_e$.
 - Given an edge $e=(i,j)$, e is in the cut if i and j are on different sides.
 - $\Pr[e \text{ in cut}] = \Pr[(i \text{ in } L) \wedge (j \text{ in } R)] + \Pr[(j \text{ in } L) \wedge (i \text{ in } R)] = 1/4 + 1/4 = 1/2$.
 - $E[X_e] = 1/2$.
 - $E[X] = e/2$ by linearity of expectations.






Correctness




- Since a cut can have at most e edges, the $e/2$ edges the algorithm outputs in expectation is a 2 approximation.
- Note that we only bounded expected size of the algorithm's cut.
 - In any particular execution, the algorithm can output a cut that's smaller or larger than $e/2$.
 - On average, the cut has size $e/2$.





Contention Resolution

An example in distributed computing where randomization is necessary

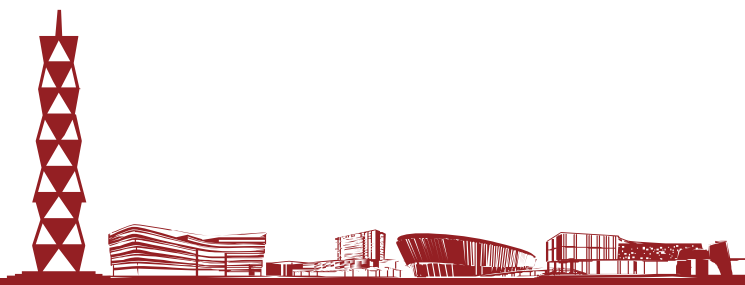




Distributed Computing



- Distributed system
 - Set of autonomous nodes, working independently of each other
 - Nodes may be able to communicate, at a cost
 - Ex: Internet, computer cluster, sensor network
- Nodes need to coordinate to solve some problem
- Coordination can be done using communication. But communication is expensive
- By making nodes randomized, they can coordinate with minimal communication
- Randomization also simplifies symmetry breaking between nodes





Contention Resolution in a Distributed System





Contention resolution. Given n processes P_1, \dots, P_n , each competing for access to a shared channel. If two or more processes access the channel simultaneously, all processes are locked out. Devise protocol to ensure all processes get through on a regular basis.

Assumption. Time is divided into rounds.

Restriction. Processes can't communicate, and they don't have id's.

Challenge. Need **symmetry-breaking** paradigm.

No deterministic protocol can solve the problem.





Contention Resolution: Randomized Protocol



Protocol. Each process requests access to the channel at time t with probability $p = 1/n$.

Claim. Let $S[i, t]$ = event that process i succeeds in accessing the database at time t . Then $1/(e \cdot n) \leq \Pr[S(i, t)] \leq 1/(2n)$.

Pf. By independence, $\Pr[S(i, t)] = p (1-p)^{n-1}$.
process i requests access none of remaining $n-1$ processes request access

- Setting $p = 1/n$, we have $\Pr[S(i, t)] = 1/n \underbrace{(1 - 1/n)^{n-1}}_{\text{between } 1/e \text{ and } 1/2}$. ▪

Useful facts from calculus.

- $1/4 < (1 - 1/n)^n < 1/e < (1 - 1/n)^{n-1} < 1/2$
- $9/4 < (1 + 1/n)^n < e < (1 + 1/n)^{n+1} < 27/8$






Contention Resolution: Randomized Protocol



Claim. The probability that process i fails to access the channel in $e \cdot n$ rounds is at most $1/e$. After $e \cdot n \cdot c \ln n$ rounds, the probability is at most n^{-c} .

Pf. Let $F[i, t]$ = event that process i fails to access database in rounds 1 through t . By independence and previous claim, we have $\Pr[F(i, t)] \leq (1 - 1/(en))^t$.

- Choose $t = e n$: $\Pr[F(i, t)] \leq \left(1 - \frac{1}{en}\right)^{en} \leq \frac{1}{e}$
- Choose $t = (e n) (c \ln n)$: $\Pr[F(i, t)] \leq \left(\frac{1}{e}\right)^{c \ln n} = n^{-c}$





Contention Resolution: Randomized Protocol



Claim. The probability that **all** processes succeed within $2e \cdot n \ln n$ rounds is at least $1 - 1/n$.

Pf. Let $F[t]$ = event that at least one of the n processes fails to access database in any of the rounds 1 through t .

$$\Pr[F[t]] = \Pr\left[\bigcup_{i=1}^n F[i,t]\right] \leq \sum_{i=1}^n \Pr[F[i,t]] \leq n \cdot n^{-c}$$

union bound previous slide

- Choosing $t = 2e n \ln n$ yields $\Pr[F[t]] \leq n \cdot n^{-2} = 1/n$. ▪

Union bound. Given events E_1, \dots, E_n , independent or not,


$$\Pr\left[\bigcup_{i=1}^n E_i\right] \leq \sum_{i=1}^n \Pr[E_i]$$





Global Minimum Cut

A problem for which the best-known randomized algorithm is faster than the best-known deterministic algorithm





Global Minimum Cut




Global min cut. Given a connected, undirected graph $G = (V, E)$, find a cut (A, B) of minimum cardinality.

Applications. Partitioning items in a database, identify clusters of related documents, network reliability, network design, circuit design, TSP solvers.

Network flow solution.

- Replace every edge (u, v) with two antiparallel edges (u, v) and (v, u) .
- Pick some vertex s and compute min s - v cut separating s from each other vertex $v \in V$.

False intuition. Global min-cut is harder than min s - t cut.







Contraction Algorithm



Contraction algorithm. [Karger 1995]

- Pick an edge $e = (u, v)$ uniformly at random.
- **Contract** edge e .
 - replace u and v by a single new supernode w
 - preserve edges, updating endpoints of u and v to w
 - keep parallel edges, but delete self-loops
- Repeat until graph has just two supernodes v_1 and v_2 .
- Return the cut (between the two supernodes).



Reference: Thore Husfeldt






Contraction Algorithm



Claim. The contraction algorithm returns a min cut with $\text{prob} \geq 2/n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G . Let F^* be edges with one endpoint in A^* and the other in B^* . Let $k = |F^*|$ = size of min cut.

- In first step, algorithm contracts an edge in F^* with $\text{prob } k / |E|$.
- Every node has degree $\geq k$ since otherwise (A^*, B^*) would not be min-cut. $\Rightarrow |E| \geq \frac{1}{2}kn$.
- Thus, algorithm contracts an edge in F^* with probability $\leq 2/n$.





Contraction Algorithm



Claim. The contraction algorithm returns a min cut with prob $\geq 2/n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G . Let F^* be edges with one endpoint in A^* and the other in B^* .
Let $k = |F^*|$ = size of min cut.

- Let G' be graph after j iterations. There are $n' = n-j$ supernodes.
 - Suppose no edge in F^* has been contracted. The min-cut in G' is still k .
 - Since value of min-cut is k , $|E'| \geq \frac{1}{2}kn' \rightarrow k/|E'| \leq 2/n'$
 - Thus, algorithm contracts an edge in F^* with probability $\leq 2/n'$.
-
- Let E_j = event that an edge in F^* is not contracted in iteration j .

$$\begin{aligned}\Pr[E_1 \cap E_2 \cdots \cap E_{n-2}] &= \Pr[E_1] \times \Pr[E_2 | E_1] \times \cdots \times \Pr[E_{n-2} | E_1 \cap E_2 \cdots \cap E_{n-3}] \\ &\geq \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \cdots \left(1 - \frac{2}{4}\right) \left(1 - \frac{2}{3}\right) \\ &= \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \cdots \left(\frac{2}{4}\right) \left(\frac{1}{3}\right) \\ &= \frac{2}{n(n-1)} \\ &\geq \frac{2}{n^2}\end{aligned}$$





Contraction Algorithm




Amplification. To amplify the probability of success, run the contraction algorithm many times.

Claim. If we repeat the contraction algorithm n^2 times with independent random choices and return the best cut found, then the algorithm finds the min-cut with constant probability.

Pf. By independence, the probability of failure is at most

$$\left(1 - \frac{2}{n^2}\right)^{n^2} = \left[\left(1 - \frac{2}{n^2}\right)^{\frac{1}{2}n^2}\right]^2 \leq (e^{-1})^2 = \frac{1}{e^2}$$






Global Min Cut: Context



Remark. Overall running time is slow since we perform $O(V^2)$ iterations, and each takes $O(E \log V)$ time (we always merge the vertex with smaller degree into the other).


Best known. [Karger 2000] $O(E \log^3 V)$.

↖
faster than best known deterministic global min cut algorithm





Random Variables and Expectations





A Quick Review of Probability Theory



Expectation. Given a discrete random variables X , its expectation $E[X]$ is defined as:


$$E[X] = \sum i \cdot \Pr[X = i]$$

Q: Roll a 6-sided dice. What is the expected value?

A: ?

Q: Roll two dice. What is the expected maximum value?

A: ?





Expectation: Two Properties



Indicator random variables. If X only takes 0 or 1, $E[X] = \Pr[X = 1]$.

Linearity of expectation. Given two random variables X and Y (not necessarily independent),

$$E[X + Y] = E[X] + E[Y].$$

Remark: $E[XY] = E[X]E[Y]$ only when X and Y are independent.

Example. Shuffle a deck of n cards; turn them over one at a time; try to guess each card. Suppose you can't remember what's been turned over already, and just guess a card from full deck uniformly at random.

Q. What's the expected number of correct guesses?

A. (surprisingly effortless using linearity of expectation)

- Let $X_i = 1$ if i^{th} guess is correct and 0 otherwise.
- Let $X = \text{number of correct guesses} = X_1 + \cdots + X_n$.
- $E[X_i] = \Pr[X_i = 1] = 1/n$.
- $E[X] = E[X_1] + \cdots + E[X_n] = 1/n + \cdots + 1/n = 1$.





Guessing Cards with Memory




Guessing with memory. Guess a card uniformly at random from cards not yet seen.

Q. What's the expected number of correct guesses?

A.

- Let $X_i = 1$ if i^{th} guess is correct and 0 otherwise.
- Let $X = \text{number of correct guesses} = X_1 + \cdots + X_n$.
- $E[X_i] = \Pr[X_i = 1] = 1/(n - i + 1)$.
- $E[X] = E[X_1] + \cdots + E[X_n] = \frac{1}{n} + \cdots + \frac{1}{2} + \frac{1}{1} = \Theta(\log n)$.





The Birthday Paradox



Problem: Suppose there are $n = 365$ days in a year, and in a room of k people, each person's birthday falls in any one of the n days with equal probability. How large should k be for us to expect two people with the same birthday?

Analysis:

- Define $X_{ij} = 1$ if person i and person j have the same birthday, and 0 otherwise.
- We know $E[X_{ij}] = \Pr[X_{ij} = 1] = 1/n$.
- Let $X = \sum_{1 \leq i < j \leq k} X_{ij}$ be the number of pairs of people having the same birthday.

- We have

$$E[X] = E\left[\sum_{1 \leq i < j \leq k} X_{ij}\right] = \binom{k}{2} \frac{1}{n} = \frac{k(k-1)}{2n}$$

- So, when $\frac{k(k-1)}{2n} \geq \frac{(k-1)^2}{2n} \geq 1$, or $k \geq \sqrt{2n} + 1 \approx 28$, we expect to see at least one pair of people having the same birthday.





Coupon Collector




Coupon collector. Each box of cereal contains a coupon. There are n different types of coupons. Assuming a box contains each type of coupon equally likely, how many boxes do you need to open to have at least one coupon of each type?

Solution.

- Stage i = time between i and $i + 1$ distinct coupons.
- Let X_i = number of steps you spend in stage i .
- Let X = number of steps in total = $X_0 + X_1 + \dots + X_{n-1}$.

$$E[X] = \sum_{i=0}^{n-1} E[X_i] = \sum_{i=0}^{n-1} \frac{n}{n-i} = n \sum_{i=1}^n \frac{1}{i} = \Theta(n \log n)$$




↑
prob of success = $(n - i)/n$
 \Rightarrow expected waiting time = $n/(n - i)$





MAX 3-SAT

An extremely simple randomized approximation algorithm





Maximum 3-Satisfiability




↙ exactly 3 distinct literals per clause

MAX-3SAT. Given 3-SAT formula, find a truth assignment that satisfies as many clauses as possible.

$$\begin{aligned}C_1 &= x_2 \vee \overline{x_3} \vee \overline{x_4} \\C_2 &= x_2 \vee x_3 \vee \overline{x_4} \\C_3 &= \overline{x_1} \vee x_2 \vee x_4 \\C_4 &= \overline{x_1} \vee \overline{x_2} \vee x_3 \\C_5 &= x_1 \vee \overline{x_2} \vee \overline{x_4}\end{aligned}$$

Remark. NP-hard search problem.

Simple idea. Flip a coin, and set each variable true with probability $\frac{1}{2}$, independently for each variable.





Maximum 3-Satisfiability: Analysis




Claim. Given a 3-SAT formula with k clauses, the **expected number** of clauses satisfied by a random assignment is $7k/8$.

Pf. Consider random variable $Z_j = \begin{cases} 1 & \text{if clause } C_j \text{ is satisfied} \\ 0 & \text{otherwise.} \end{cases}$

- Let Z = total number of clauses satisfied.

$$\begin{aligned} E[Z] &= \sum_{j=1}^k E[Z_j] \\ \text{linearity of expectation} &= \sum_{j=1}^k \Pr[\text{clause } C_j \text{ is satisfied}] \\ &= \frac{7}{8}k \end{aligned}$$





Maximum 3-Satisfiability: Analysis

Lemma. The probability that a random assignment satisfies $\geq 7k/8$ clauses is at least $1/(8k)$.

Pf. Let p_j be probability that exactly j clauses are satisfied.

We start by writing

$$\begin{aligned}\frac{7}{8}k &= \sum_{j=0}^k jp_j = \sum_{j < 7k/8} jp_j + \sum_{j \geq 7k/8} jp_j \\ &\leq \sum_{j < 7k/8} k'p_j + \sum_{j \geq 7k/8} kp_j \\ &= k'(1-p) + kp \leq k' + kp\end{aligned}$$

Hence, $kp \geq \frac{7}{8}k - k'$

But $\frac{7}{8}k - k' \geq 1/8$ (k' is the largest natural number that is strictly smaller than $\frac{7}{8}k$)

So

$$p \geq \frac{\frac{7}{8}k - k'}{k} \geq \frac{1}{8k}.$$





Next Time: Randomized algorithms (Cont.)

