

# Announcement

---

- ▶ Homework 3
  - ▶ Available in Blackboard -> Homework
  - ▶ Due: Apr. 18, 11:59pm



## The rest of the course...


---

- ▶ We will cover more traditional NLP tasks & techniques.
- ▶ They are the past.
- ▶ They may also be (part of) the future.





# Sequence Labeling



SLP3 Ch 8, 9.4; INLP Ch 7, 8

# Sequence Labeling

---

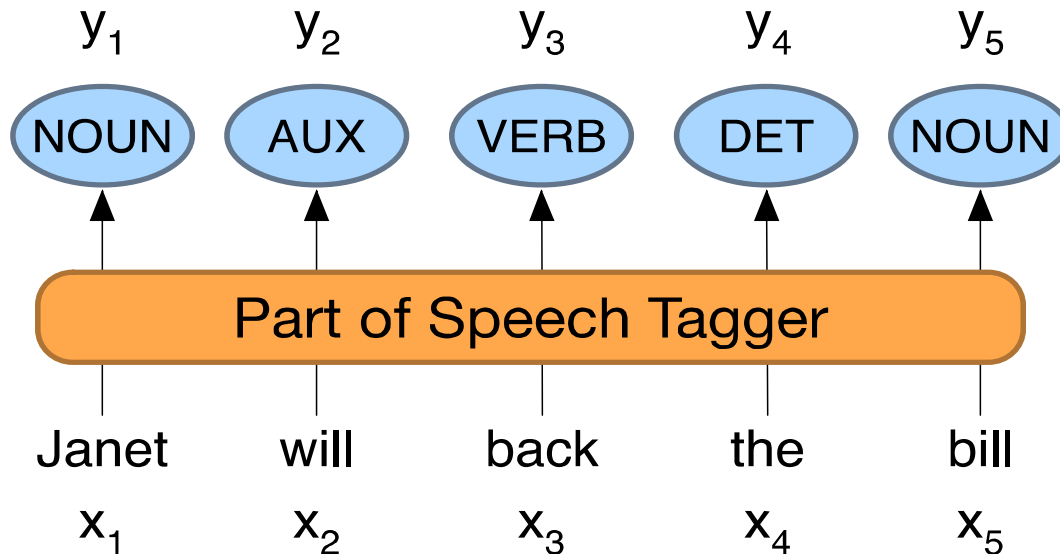
- ▶ Known
  - ▶ A set of labels  $Y = \{y^1, y^2, \dots, y^n\}$
- ▶ Input:
  - ▶ Sentence  $x = \{x_1, x_2, \dots, x_m\}$
- ▶ Output:
  - ▶ For each word  $x_i$ , predict a label  $y_i \in Y$



# Part-of-Speech (POS) Tagging

---

- ▶ Map from sequence  $x_1, x_2, \dots, x_m$  of words to  $y_1, y_2, \dots, y_m$  of POS tags



# Two classes of words: Open vs. Closed

---

## ▶ Closed class words

- ▶ Relatively fixed for each language
- ▶ Usually **function** words: short, frequent words with grammatical function
  - ▶ determiners: *a, an, the*
  - ▶ pronouns: *she, he, I*
  - ▶ prepositions: *on, under, over, near, by, ...*

## ▶ Open class words

- ▶ Usually **content** words: Nouns, Verbs, Adjectives, Adverbs
  - ▶ Plus interjections: *oh, ouch, uh-huh, yes, hello*



# Two classes of words: Open vs. Closed

---

## Open class ("content") words

### Nouns

#### Proper

*Janet*  
*Italy*

#### Common

*cat, cats*  
*mango*

### Verbs

#### Main

*eat*  
*went*

### Adjectives

*old green tasty*

### Adverbs

*slowly yesterday*

### Numbers

*122,312*  
*one*

Interjections *Ow hello*

*... more*

## Closed class ("function")

Determiners *the some*

Conjunctions *and or*

Pronouns *they its*

### Auxiliary

*can*  
*had*

Prepositions *to with*

Particles *off up*

*... more*



# Why Part-of-Speech Tagging

---

- ▶ Can be useful for other NLP tasks
  - ▶ Parsing
    - ▶ POS tagging can improve syntactic parsing
  - ▶ MT
    - ▶ reordering of adjectives and nouns (say from Spanish to English)
  - ▶ Sentiment or affective tasks
    - ▶ may want to distinguish adjectives or other POS
  - ▶ Text-to-speech
    - ▶ how do we pronounce “lead” or “object”?





# Other sequence labeling tasks

---

## ▶ Chinese word segmentation

### ▶ Input

瓦 里 西 斯 的 船 只 中 ...

### ▶ Output

B I I E S B E S ...

(瓦 里 西 斯) (的) (船 只) (中) ...

B = beginning of a word

I = inside of a word

E = end of a word

S = single character word



# Other sequence labeling tasks

---

## ▶ Named entity recognition

### ▶ Input

Michael Jeffrey Jordan was born in Brooklyn ...

### ▶ Output

|                               |       |       |                 |   |   |       |
|-------------------------------|-------|-------|-----------------|---|---|-------|
| B-PER                         | I-PER | E-PER | O               | O | O | S-LOC |
| <u>Michael Jeffrey Jordan</u> |       |       | <u>Brooklyn</u> |   |   |       |
| Person                        |       |       | Location        |   |   |       |

B = beginning of an entity

I = inside of an entity

E = end of an entity

S = single word entity

O = outside of any entity

-PER = person

-LOC = location

-ORG = organization

...



# Other sequence labeling tasks

---

## ▶ Semantic role labeling

### ▶ Input

The cat loves hats ...

### ▶ Output

B-ARG0 E-ARG0 S-PRED S-ARG1  
The cat ← arg0 loves → hats arg1

B = beginning of an entity

I = inside of an entity

E = end of an entity

S = single word entity

O = outside of any entity

-PRED = predicate

-ARG0 = agent

-ARG1 = patient

...



# The simplest method

---

- ▶ For each word, predict its most frequent label
  - ▶ 90% accuracy on POS tagging!
  - ▶ Disadvantages:
    1. It does not consider the contextual info
      - ▶ “book a flight” vs. “read a book”
      - ▶ 我骑车差点摔倒，好在我一把把把把住了
      - ▶ 校长说衣服上除了校徽别别别的
    2. It does not consider relations between adjacent labels
      - ▶ In BIOES: "B-I" and "B-E" are OK, but "B-O" and "B-S" are not



# Method

---

- ▶ Hidden Markov model (HMM)
- ▶ Conditional random field (CRF)
- ▶ Neural models





# Hidden Markov Model



# Hidden Markov Model (HMM)

---

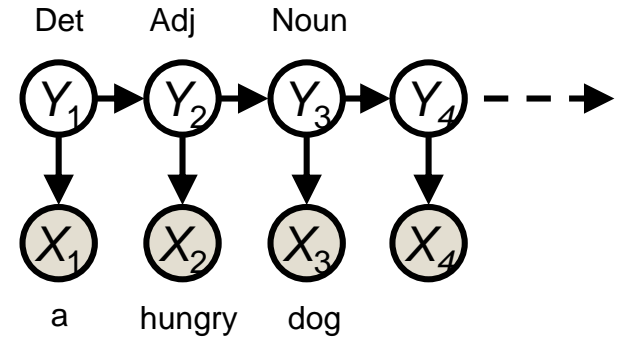
## ▶ Variables

- ▶  $X$ : word
- ▶  $Y$ : label (hidden state)

## ▶ Parameters

- ▶ Transition model  $P(y_t|y_{t-1})$ 
  - ▶ Similar to a bigram model
- ▶ Emission model  $P(x_t|y_t)$
- ▶ Initial distribution  $P(y_1)$ 
  - ▶ Can be seen as transition from  $Y_0=\text{START}$  to  $Y_1$
- ▶ Modeling end of sequence
  - ▶ Can be seen as transition from  $Y_n$  to  $Y_{n+1}=\text{STOP}$

- ▶ Joint prob:  $P(x_1 \cdots x_n, y_1 \cdots y_{n+1}) = \prod_t P(y_t|y_{t-1})P(x_t|y_t)$



# HMM Example


---

**Transition**

| $Y_{t-1}$ | $P(Y_t Y_{t-1})$ |     |     |     |
|-----------|------------------|-----|-----|-----|
|           | N                | V   | P   | ... |
| START     | 0.5              | 0.1 | 0.1 | ... |
| N         | 0.4              | 0.3 | 0.1 | ... |
| V         | 0.5              | 0   | 0.3 | ... |
| P         | 0.3              | 0.1 | 0   | ... |
| ...       | ...              | ... | ... | ... |

**Emission**

| $Y_t$ | $P(X_t Y_t)$ |       |       |     |
|-------|--------------|-------|-------|-----|
|       | "fox"        | "dog" | "run" | ... |
| N     | 0.02         | 0.03  | 0.01  | ... |
| V     | 0            | 0     | 0.05  | ... |
| P     | 0            | 0     | 0     | ... |
| ...   | ...          | ...   | ...   | ... |






# High-order HMM

---

- ▶ Transition model  $P(y_t | y_{t-1}, y_{t-2}, \dots, y_{t-n+1})$ 
  - ▶ Similar to an n-gram model



# HMM Inference (Decoding)

---

- ▶ Given an input sequence, find the most likely label sequence under the model

$$y^* = \operatorname{argmax}_{y_1 \cdots y_n} P(y_1 \cdots y_{n+1} | x_1 \cdots x_n) = \operatorname{argmax}_{y_1 \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1})$$

- ▶ Given an input, we can score any tag sequence

- ▶  $P(x_1 \cdots x_n, y_1 \cdots y_{n+1}) = \prod_t P(y_t | y_{t-1}) P(x_t | y_t)$

|     |        |          |       |     |         |   |
|-----|--------|----------|-------|-----|---------|---|
| NNP | VBZ    | NN       | NNS   | CD  | NN      | . |
| Fed | raises | interest | rates | 0.5 | percent | . |

$q(\text{NNP}|\text{START}) e(\text{Fed}|\text{NNP}) q(\text{VBZ}|\text{NNP}) e(\text{raises}|\text{VBZ}) q(\text{NN}|\text{VBZ}) \dots$



# HMM Inference (Decoding)

---

- ▶ Given an input sequence, find the most likely label sequence under the model

$$y^* = \operatorname{argmax}_{y_1 \cdots y_n} P(y_1 \cdots y_{n+1} | x_1 \cdots x_n) = \operatorname{argmax}_{y_1 \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1})$$

- ▶ Given an input, we can score any tag sequence
- ▶ In principle, we're done – list all possible tag sequences, score each one, pick the best one
  - ▶ Exponential time complexity!

|                      |   |            |
|----------------------|---|------------|
| NNP VBZ NN NNS CD NN | ➡ | logP = -23 |
|----------------------|---|------------|

|                      |   |            |
|----------------------|---|------------|
| NNP NNS NN NNS CD NN | ➡ | logP = -29 |
|----------------------|---|------------|

|                      |   |            |
|----------------------|---|------------|
| NNP VBZ VB NNS CD NN | ➡ | logP = -27 |
|----------------------|---|------------|

... ..

... ..



# Dynamic Programming (Viterbi Algorithm)

---

- ▶ Define  $\pi(i, y_i)$  to be the max score of a tag sequence of length  $i$  ending in tag  $y_i$

$$\begin{aligned}\pi(i, y_i) &= \max_{y_1 \cdots y_{i-1}} P(x_1 \cdots x_i, y_1 \cdots y_i) \\&= \max_{y_1 \cdots y_{i-1}} e(x_i | y_i) q(y_i | y_{i-1}) P(x_1 \cdots x_{i-1}, y_1 \cdots y_{i-1}) \\&= e(x_i | y_i) \max_{y_{i-1}} q(y_i | y_{i-1}) \max_{y_1 \cdots y_{i-2}} P(x_1 \cdots x_{i-1}, y_1 \cdots y_{i-1}) \\&= e(x_i | y_i) \max_{y_{i-1}} q(y_i | y_{i-1}) \pi(i-1, y_{i-1})\end{aligned}$$



# Dynamic Programming (Viterbi Algorithm)

---

- ▶ Define  $\pi(i, y_i)$  to be the max score of a tag sequence of length  $i$  ending in tag  $y_i$

$$\pi(i, y_i) = e(x_i | y_i) \max_{y_{i-1}} q(y_i | y_{i-1}) \pi(i-1, y_{i-1})$$

- ▶ We now have an efficient DP algorithm
  - ▶ Start with  $\pi(0, \text{START}) = 1$
  - ▶ Work your way to the end of the sentence

$$\begin{aligned} P(y^*) &= \max_{y_1 \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1}) \\ &= \max_{y_n} q(\text{STOP} | y_n) \max_{y_1 \cdots y_{n-1}} P(x_1 \cdots x_n, y_1 \cdots y_n) \\ &= \max_{y_n} q(\text{STOP} | y_n) \pi(n, y_n) := \pi(n+1, \text{STOP}) \end{aligned}$$



# Example

## State Trellis

Fruit

Flies

Like

Bananas

$\pi(1, N)$

$\pi(2, N)$

$\pi(3, N)$

$\pi(4, N)$

$\pi(1, V)$

$\pi(2, V)$

$\pi(3, V)$

$\pi(4, V)$

$\pi(1, IN)$

$\pi(2, IN)$

$\pi(3, IN)$

$\pi(4, IN)$

START

STOP

$\pi(0, \text{START})$   
 $= 1$

$$\pi(i, y_i) = e(x_i | y_i) \max_{y_{i-1}} q(y_i | y_{i-1}) \pi(i-1, y_{i-1})$$



# Example

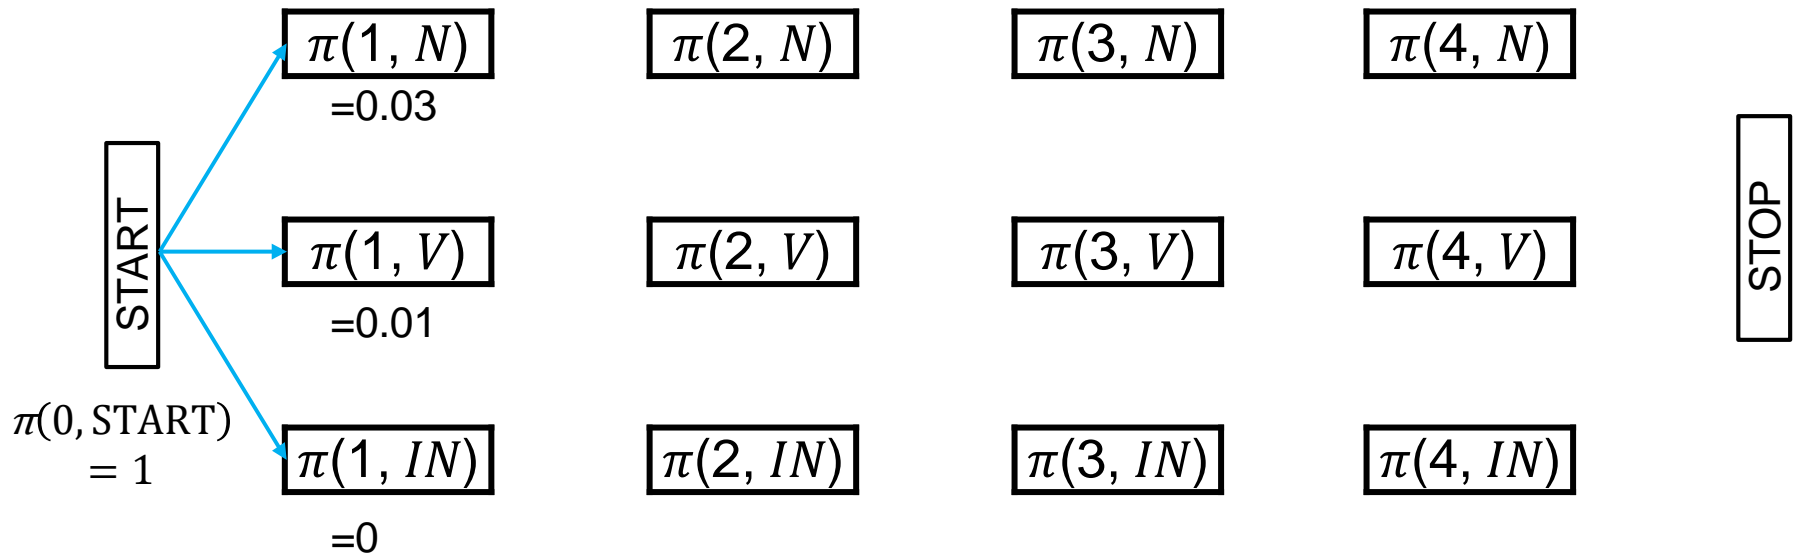
---

Fruit

Flies

Like

Bananas



$$\pi(i, y_i) = e(x_i | y_i) \max_{y_{i-1}} q(y_i | y_{i-1}) \pi(i-1, y_{i-1})$$




# Example

Fruit

Flies

Like

Bananas



$$\pi(i, y_i) = e(x_i | y_i) \max_{y_{i-1}} q(y_i | y_{i-1}) \pi(i-1, y_{i-1})$$



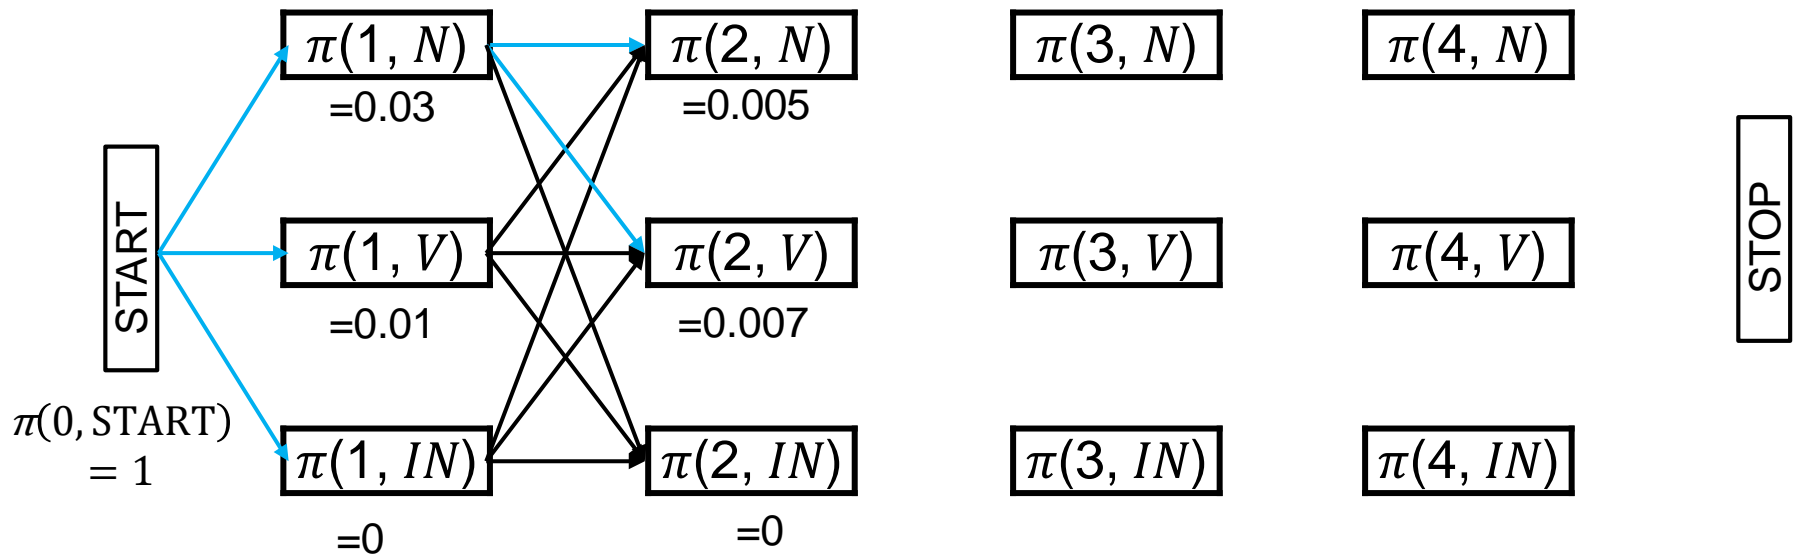
# Example

Fruit

Flies

Like

Bananas



$$\pi(i, y_i) = e(x_i | y_i) \max_{y_{i-1}} q(y_i | y_{i-1}) \pi(i-1, y_{i-1})$$


# Example

Fruit

Flies

Like

Bananas



$$\pi(i, y_i) = e(x_i | y_i) \max_{y_{i-1}} q(y_i | y_{i-1}) \pi(i-1, y_{i-1})$$


# Example

Fruit

Flies

Like

Bananas



$$\pi(i, y_i) = e(x_i | y_i) \max_{y_{i-1}} q(y_i | y_{i-1}) \pi(i-1, y_{i-1})$$


# Example

Fruit

Flies


Like

Bananas



$$\pi(i, y_i) = e(x_i | y_i) \max_{y_{i-1}} q(y_i | y_{i-1}) \pi(i-1, y_{i-1})$$

# Bananas



$$\pi(i, y_i) = e(x_i | y_i) \max_{y_{i-1}} q(y_i | y_{i-1}) \pi(i-1, y_{i-1})$$



# The Viterbi Algorithm: Runtime

---

$$\pi(i, y_i) = e(x_i | y_i) \max_{y_{i-1}} q(y_i | y_{i-1}) \pi(i-1, y_{i-1})$$

- ▶ Sentence length  $n$ , tag number  $|Y|$
- ▶  $O(n |Y|)$  entries in  $\pi(i, y_i)$
- ▶  $O(|Y|)$  time to compute each  $\pi(i, y_i)$
- ▶ Total runtime:  $O(n |Y|^2)$



# Marginal Inference

---

- ▶ Compute the marginal probability of the input sentence
  - ▶  $P(x_1 \cdots x_n) = \sum_{y_1 \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1})$
- ▶ Given an input, we can score any tag sequence
- ▶ In principle, we're done – list all possible tag sequences with  $y_i$ , score each one, take summation
  - ▶ Exponential time complexity!

|                      |   |            |
|----------------------|---|------------|
| NNP VBZ NN NNS CD NN | ➡ | logP = -23 |
| NNP NNS NN NNS CD NN | ➡ | logP = -29 |
| NNP VBZ VB NNS CD NN | ➡ | logP = -27 |
| ... ..               |   | ... ..     |



# Marginal Inference

---

- ▶ Compute the marginal probability of the input sentence

- ▶  $P(x_1 \cdots x_n) = \sum_{y_1 \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1})$

- ▶ Compare it with decoding

- ▶  $y^* = \operatorname{argmax}_{y_1 \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1})$


- ▶  $P(y^*) = \max_{y_1 \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1})$





# The State Trellis: Viterbi


$$\max_{y_1 \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1})$$



# The State Trellis: Marginal

---

$$\sum_{y_1 \dots y_n} P(x_1 \dots x_n, y_1 \dots y_{n+1})$$



# Dynamic Programming (Forward Algorithm)

---

$$\begin{aligned}\alpha(i, y_i) &= P(x_1 \cdots x_i, y_i) = \sum_{y_1, \dots, y_{i-1}} P(x_1 \cdots x_i, y_1 \cdots y_i) \\&= \sum_{y_1, \dots, y_{i-1}} e(x_i | y_i) q(y_i | y_{i-1}) P(x_1 \cdots x_{i-1}, y_1 \cdots y_{i-1}) \\&= \sum_{y_{i-1}} e(x_i | y_i) q(y_i | y_{i-1}) \sum_{y_1, \dots, y_{i-2}} P(x_1 \cdots x_{i-1}, y_1 \cdots y_{i-1}) \\&= \sum_{y_{i-1}} e(x_i | y_i) q(y_i | y_{i-1}) \alpha(i-1, y_{i-1})\end{aligned}$$



# Dynamic Programming (Forward Algorithm)

---

- ▶ Start with:

$$\alpha(0, y_0) = \begin{cases} 1 & \text{if } y_0 = \text{START} \\ 0 & \text{otherwise} \end{cases}$$

- ▶ For  $i = 1, \dots, n$ :

$$\alpha(i, y_i) = \sum_{y_{i-1}} e(x_i | y_i) q(y_i | y_{i-1}) \alpha(i-1, y_{i-1})$$

- ▶ Finally:

$$\begin{aligned} P(x_1 \cdots x_n) &= \sum_{y_1 \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1}) \\ &= \sum_{y_n} q(\text{STOP} | y_n) \sum_{y_1 \cdots y_{n-1}} P(x_1 \cdots x_n, y_1 \cdots y_n) \\ &= \sum_{y_n} q(\text{STOP} | y_n) \alpha(n, y_n) \coloneqq \alpha(n+1, \text{STOP}) \end{aligned}$$



# Marginal Inference

---


- ▶ Find the marginal probability of each tag for  $y_i$ 
  - ▶  $P(x_1 \cdots x_n, y_i) = \sum_{y_1 \cdots y_{i-1}} \sum_{y_{i+1} \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1})$
- ▶ Given an input, we can score any tag sequence
- ▶ In principle, we're done – list all possible tag sequences with  $y_i$ , score each one, take summation
  - ▶ Exponential time complexity!

|     |     |    |     |    |    |   |            |
|-----|-----|----|-----|----|----|---|------------|
| NNP | VBZ | NN | NNS | CD | NN | ⇒ | logP = -23 |
| NNP | NNS | NN | NNS | CD | NN | ⇒ | logP = -29 |
| NNP | VBZ | VB | NNS | CD | NN | ⇒ | logP = -27 |
| ... | ... |    |     |    |    |   | ...        |



# The State Trellis: Marginal

$$\sum_{y_1 \dots y_{i-1}} \sum_{y_{i+1} \dots y_n} P(x_1 \dots x_n, y_1 \dots y_{n+1})$$




# Dynamic Programming

---

$$P(x_1 \cdots x_n, y_i) = \sum_{y_1 \cdots y_{i-1}} \sum_{y_{i+1} \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1})$$


$$\begin{aligned} P(x_1 \cdots x_n, y_i) &= P(x_1 \cdots x_i, y_i) P(x_{i+1} \cdots x_n | y_i, x_1 \cdots x_i) \\ &= \underbrace{P(x_1 \cdots x_i, y_i)}_{\alpha(i, y_i)} \underbrace{P(x_{i+1} \cdots x_n | y_i)}_{\beta(i, y_i)} \end{aligned}$$



# Forward

---

$$\begin{aligned}\alpha(i, y_i) &= P(x_1 \cdots x_i, y_i) = \sum_{y_1, \dots, y_{i-1}} P(x_1 \cdots x_i, y_1 \cdots y_i) \\&= \sum_{y_1, \dots, y_{i-1}} e(x_i | y_i) q(y_i | y_{i-1}) P(x_1 \cdots x_{i-1}, y_1 \cdots y_{i-1}) \\&= \sum_{y_{i-1}} e(x_i | y_i) q(y_i | y_{i-1}) \sum_{y_1, \dots, y_{i-2}} P(x_1 \cdots x_{i-1}, y_1 \cdots y_{i-1}) \\&= \sum_{y_{i-1}} e(x_i | y_i) q(y_i | y_{i-1}) \alpha(i-1, y_{i-1})\end{aligned}$$





# Backward


---

$$\beta(i, y_i) = P(x_{i+1} \cdots x_n | y_i) = \sum_{y_{i+1}, \dots, y_n} P(x_{i+1} \cdots x_n, y_{i+1} \cdots y_{n+1} | y_i)$$

$$= \sum_{y_{i+1}, \dots, y_n} e(x_{i+1} | y_{i+1}) q(y_{i+1} | y_i) P(x_{i+2} \cdots x_n, y_{i+2} \cdots y_{n+1} | y_{i+1})$$

$$= \sum_{y_{i+1}} e(x_{i+1} | y_{i+1}) q(y_{i+1} | y_i) \sum_{y_{i+2}, \dots, y_n} P(x_{i+2} \cdots x_n, y_{i+2} \cdots y_{n+1} | y_{i+1})$$

$$= \sum_{y_{i+1}} e(x_{i+1} | y_{i+1}) q(y_{i+1} | y_i) \beta(i + 1, y_{i+1})$$



# Forward-Backward Algorithm

---

- ▶ Two passes: one forward, one backward

- ▶ Forward

$$\alpha(0, y_0) = \begin{cases} 1 & \text{if } y_0 = \text{START} \\ 0 & \text{otherwise} \end{cases}$$

- ▶ For  $i = 1, \dots, n$ :

$$\alpha(i, y_i) = \sum_{y_{i-1}} e(x_i | y_i) q(y_i | y_{i-1}) \alpha(i-1, y_{i-1})$$

- ▶ Backward

$$\beta(n, y_n) = q(\text{STOP} | y_n)$$

- ▶ For  $i = n-1, \dots, 1$

$$\beta(i, y_i) = \sum_{y_{i+1}} e(x_{i+1} | y_{i+1}) q(y_{i+1} | y_i) \beta(i+1, y_{i+1})$$



# Forward-Backward: Runtime

---

$$\alpha(i, y_i) = \sum_{y_{i-1}} e(x_i | y_i) q(y_i | y_{i-1}) \alpha(i-1, y_{i-1})$$

$$\beta(i, y_i) = \sum_{y_{i+1}} e(x_{i+1} | y_{i+1}) q(y_{i+1} | y_i) \beta(i+1, y_{i+1})$$

- ▶ Sentence length  $n$ , tag number  $|Y|$
- ▶  $O(n|Y|)$  entries in  $\alpha(i, y_i)$  and  $\beta(i, y_i)$
- ▶  $O(|Y|)$  time to compute each entry
- ▶ Total runtime:  $O(n|Y|^2)$
- ▶ Exactly the same as Viterbi



# HMM Supervised Learning

---

- ▶ Learn HMM given annotated sequence  $\{(x_1, y_1), \dots, (x_n, y_n)\}$ 
  - ▶ Maximum likelihood estimate

$$P(x, y) = \prod_{i=1}^{n+1} e(x_i | y_i) \cdot q(y_i | y_{i-1}) = \prod_{i,j \in Y} q(j|i)^{c(i,j)} \prod_{j \in X} \prod_{i \in Y} e(j|i)^{c(i,j)}$$

$e$ : emission;  $q$ : transition;  $c$ : co-occurrence count

- ▶ Closed-form solution: count and normalize

$$e(k|i) = \frac{c(i,k)}{\sum_{k' \in X} c(i,k')} \quad q(j|i) = \frac{c(i,j)}{\sum_{j' \in Y} c(i,j')}$$

- ▶ Handle data sparseness
  - ▶ We can use all of the tricks we use for n-gram models



# HMM Unsupervised Learning

---

- ▶ Learn HMM given **unannotated** sequence  $\{x_1, \dots, x_n\}$
- ▶ Application: part-of-speech induction
  - ▶ Induce the set of POS tags from text
- ▶ Maximize marginal likelihood

$$P(x_1 \cdots x_n) = \sum_{y_1 \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1})$$



# Expectation-Maximization (EM)

---

- ▶ Can be used to learn any model with hidden variables (missing data)
- ▶ Alternate:
  - ▶ Compute distributions over hidden variables based on current parameter values
  - ▶ Compute new parameter values to maximize expected log likelihood based on distributions over hidden variables
- ▶ Stop when no changes
- ▶ Can reach a local optimum but not necessarily a global optimum



# EM for HMM (Baum-Welch Algorithm )

---

- ▶ Initialize transition and emission parameters
  - ▶ Random, uniform, or more informed initialization

- ▶ Iterate until convergence

- ▶ E-Step:

- ▶ Compute expected counts

- ▶ General form:

$$c(S) = E_{P(y_1 \cdots y_{n+1} | x_1 \cdots x_n)} [Count(S | x_1, \cdots, x_n, y_1 \cdots y_{n+1})]$$

These statistics summarize  
 $P(y_1 \cdots y_{n+1} | x_1 \cdots x_n)$

- ▶ Computing:

$$c(NN) = \sum_i c(y_i = NN) = \sum_i P(y_i = NN | x_1, \cdots, x_n)$$

$$c(NN \rightarrow VB) = \cdots = \sum_i P(y_i = NN, y_{i+1} = VB | x_1, \cdots, x_n)$$

$$c(NN \rightarrow apple) = \cdots = \sum_i P(y_i = NN, x_i = apple | x_1, \cdots, x_n)$$

- ▶ These are for one sentence. Take sum if multiple sentences.
- 



# Compute expected counts

---

$$\begin{aligned}c(NN) &= \sum_i P(y_i = NN | x_1, \dots, x_n) \\&= \sum_i \frac{P(x_1 \dots x_n, y_i = NN)}{P(x_1 \dots x_n)} \\&= \frac{\sum_i \alpha(i, y_i = NN) \beta(i, y_i = NN)}{\alpha(n+1, STOP)}\end{aligned}$$

$$\begin{aligned}c(NN \rightarrow VB) &= \sum_i P(y_i = NN, y_{i+1} = VB | x_1, \dots, x_n) \\&= \sum_i \frac{P(x_1 \dots x_n, y_i = NN, y_{i+1} = VB)}{P(x_1 \dots x_n)} \\&= \frac{\sum_i \alpha(i, y_i = NN) q(VB|NN) e(x_{i+1}|VB) \beta(i+1, y_{i+1} = VB)}{\alpha(n+1, STOP)}\end{aligned}$$





# EM for HMM (Baum-Welch Algorithm )

---

- ▶ Initialize transition and emission parameters
  - ▶ Random, uniform, or more informed initialization
- ▶ Iterate until convergence
  - ▶ E-Step:
    - ▶ Compute expected counts
  - ▶ M-step:
    - ▶ Compute new parameter values to maximize expected log likelihood

These statistics summarize  
 $P(y_1 \cdots y_{n+1} | x_1 \cdots x_n)$

$$E_{Q(y_1 \cdots y_{n+1})}[\log P(x_1, \cdots, x_n, y_1 \cdots y_{n+1})]$$

- ▶ Closed form solution: normalizing expected counts

$$e_{ML}(x|y) = \frac{c(y, x)}{c(y)} \qquad q_{ML}(y_i|y_{i-1}) = \frac{c(y_{i-1}, y_i)}{c(y_{i-1})}$$



# HMM Unsupervised Learning

---

- ▶ Learn HMM given **unannotated** sequence  $\{x_1, \dots, x_n\}$
- ▶ Maximize marginal likelihood

$$P(x_1 \cdots x_n) = \sum_{y_1 \cdots y_n} P(x_1 \cdots x_n, y_1 \cdots y_{n+1})$$

- ▶ EM for HMM (Baum-Welch Algorithm )
- ▶ Can we directly optimize it by gradient descent?
  - ▶ Yes!
  - ▶ Use forward to compute  $P(x_1 \cdots x_n)$
  - ▶ Run backprop on the computation graph



# Forward-Backward is just backprop!

---

- ▶ The forward and then backprop procedure is almost the same as Forward-Backward
- ▶ Expected counts can be computed by backprop
  - ▶  $c(NN \rightarrow VB) = \frac{\partial \log P(x_1 \cdots x_n)}{\partial q(NN \rightarrow VB)}$
  - ▶  $c(NN \rightarrow apples) = \frac{\partial \log P(x_1 \cdots x_n)}{\partial e(NN \rightarrow apples)}$
- ▶ See <https://aclanthology.org/W16-5901.pdf>






## From HMM to Conditional Random Field



# Beyond HMM

---

- ▶ The simplest method: for each word, predict its most frequent label
  - ▶ Problems:
    - ☹️ 1. It does not consider the contextual info
    - 😊 2. It does not consider relations between adjacent labels
- ▶ Does HMM solve the two problems?
  - ▶ HMM handles problem 2, but not 1



# Max-Entropy Markov Models (MEMM)

---


- ▶ From HMM to MEMM

- ▶ HMM is a generative model

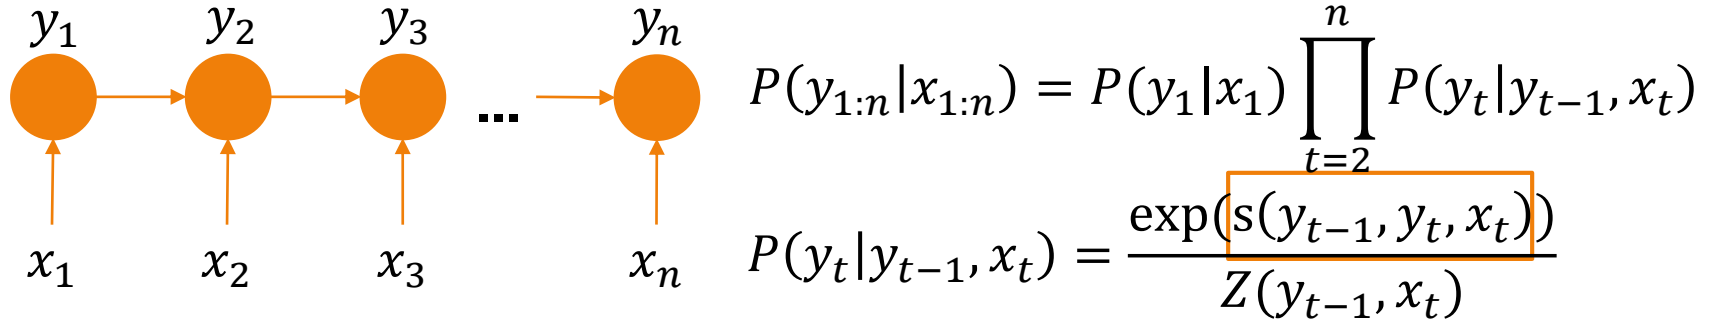
$$P(x_{1:n}, y_{1:n}) = \prod_t P(y_t | y_{t-1}) P(x_t | y_t)$$

- ▶ MEMM is a discriminative model

$$P(y_{1:n} | x_{1:n}) = P(y_1 | x_1) \prod_{t=2}^n P(y_t | y_{t-1}, x_t)$$



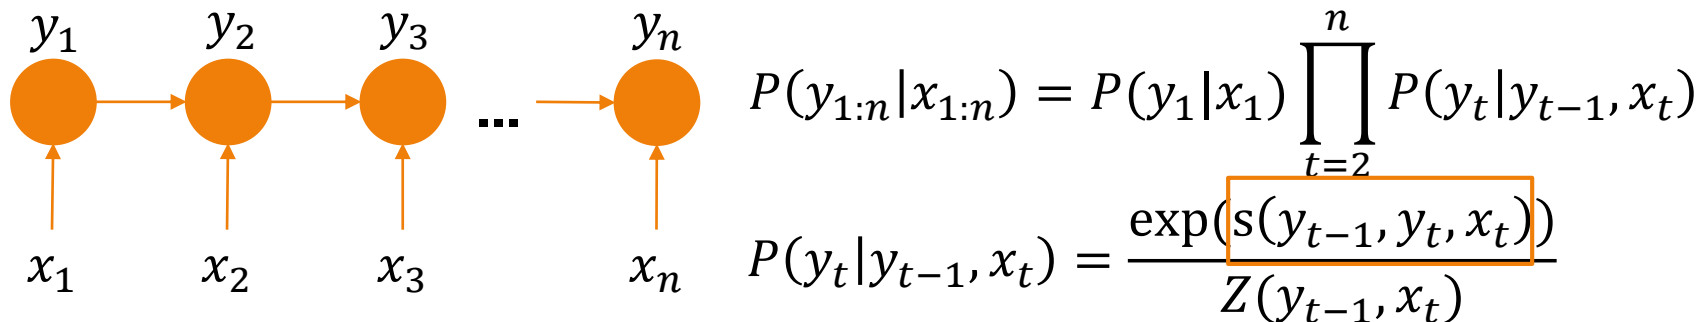
# Max-Entropy Markov Models (MEMM)



- ▶ Score function  $s(y_{t-1}, y_t, x_t)$  can be a simple linear function  $W^T f(y_{t-1}, y_t, x_t)$ .
  - ▶ Possible features:
    - ▶  $y_{t-1}$  is B and  $y_t$  is E?
    - ▶  $y_{t-1}$  is B and  $y_t$  is O?
    - ▶  $x_t$  is a noun?
    - ▶  $x_t$  is capitalized? ...



# Max-Entropy Markov Models (MEMM)




- ▶ Score function  $s(y_{t-1}, y_t, x_t)$  can be a simple linear function  $W^T f(y_{t-1}, y_t, x_t)$ .
- ▶ It may also be a neural network with word embedding of  $x_t$  and label embedding of  $y_t$  and  $y_{t-1}$  as input
  - ▶ more on this later...
- ▶ Sometimes,  $s(y_{t-1}, y_t, x_t)$  is decomposed to a transition score and an emission score
  - ▶  $s(y_{t-1}, y_t, x_t) = s_e(y_t, x_t) + s_q(y_t, y_{t-1})$




# Max-Entropy Markov Models (MEMM)

---



$$P(y_{1:n}|x_{1:n}) = P(y_1|x_1) \prod_{t=2}^n P(y_t|y_{t-1}, x_t)$$

$$P(y_t|y_{t-1}, x_t) = \frac{\exp(s(y_{t-1}, y_t, x_t))}{Z(y_{t-1}, x_t)}$$



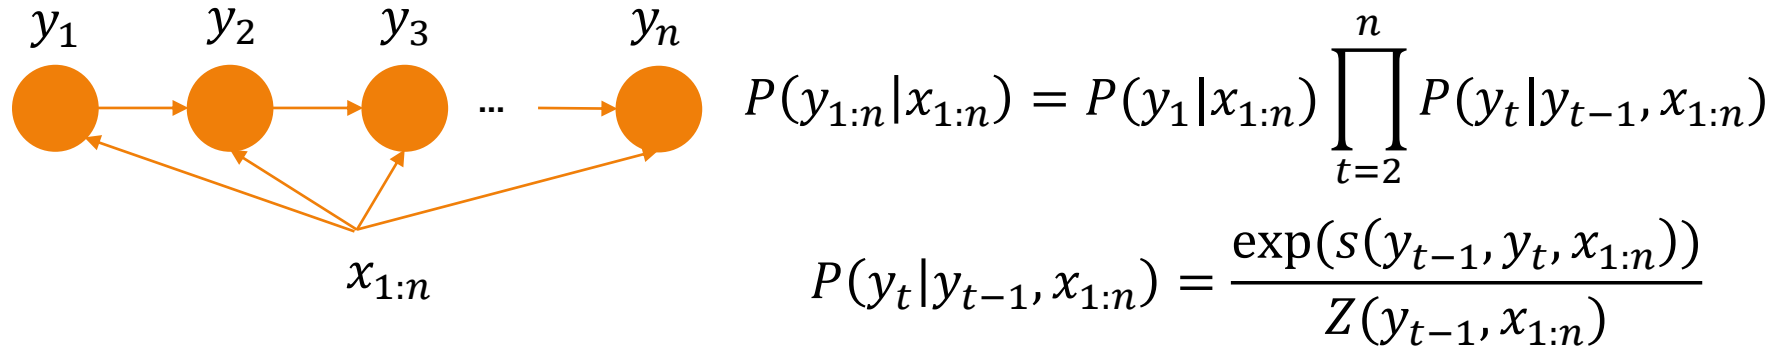
$$P(y_{1:n}|x_{1:n}) = P(y_1|x_{1:n}) \prod_{t=2}^n P(y_t|y_{t-1}, x_{1:n})$$

$$P(y_t|y_{t-1}, x_{1:n}) = \frac{\exp(s(y_{t-1}, y_t, x_{1:n}))}{Z(y_{t-1}, x_{1:n})}$$



# Max-Entropy Markov Models (MEMM)

---




- ▶ Now we can consider info from the whole sentence in the score function
- ▶ MEMM considers both contextual info and relations between adjacent labels!
- ▶ But MEMM suffers from label bias problem:
  - ▶ Preference of states with lower number of transitions, because transitions from such a state have higher average probabilities.




# From MEMM to CRF

---



$$P(y_{1:n}|x_{1:n}) = P(y_1|x_{1:n}) \prod_{t=2}^n P(y_t|y_{t-1}, x_{1:n})$$

$$P(y_t|y_{t-1}, x_{1:n}) = \frac{\exp(s(y_{t-1}, y_t, x_{1:n}))}{Z(y_{t-1}, x_{1:n})}$$




$$P(y_{1:n}|x_{1:n}) = \frac{1}{Z(x_{1:n})} \prod_{t=1}^n \exp(s(y_{t-1}, y_t, x_{1:n}))$$



# From MEMM to CRF

---



$$P(y_{1:n}|x_{1:n}) = \frac{1}{Z(x_{1:n})} \prod_{t=1}^n \exp(s(y_{t-1}, y_t, x_{1:n}))$$

- ▶ Conditional Random Field (CRF) is an undirected graphical model
  - ▶ Global normalization instead of local normalization
  - ▶ Both problems solved ✓
  - ▶ Label bias solved ✓
    - ▶ States with lower number of transitions can still have low scores.



# CRF inference (decoding)

$$y^* = \operatorname{argmax}_{y_1 \cdots y_n} \frac{1}{Z(x_{1:n})} \prod_{t=1}^n \exp(s(y_{t-1}, y_t, x_{1:n}))$$

$$= \operatorname{argmax}_{y_1 \cdots y_n} \prod_{t=1}^n \exp(s(y_{t-1}, y_t, x_{1:n}))$$

$$= \operatorname{argmax}_{y_1 \cdots y_n} \sum_{t=1}^n s(y_{t-1}, y_t, x_{1:n})$$

Score of label sequence  
 $s(y_{1:n})$

## ► Decoding by Viterbi

$$\pi(i, y_i) = \max_{y_1 \cdots y_{i-1}} \sum_{t=1}^i s(y_{t-1}, y_t, x_{1:n})$$

$$\pi(0, \text{START}) = 0$$

$$= \max_{y_{i-1}} s(y_{i-1}, y_i, x_{1:n}) + \max_{y_1 \cdots y_{i-2}} \sum_{t=1}^{i-1} s(y_{t-1}, y_t, x_{1:n})$$

$$= \max_{y_{i-1}} s(y_{i-1}, y_i, x_{1:n}) + \pi(i-1, y_{i-1})$$



# CRF Supervised Learning

---

- ▶ Learn CRF given annotated sequence  $\{(x_1, y_1), \dots, (x_n, y_n)\}$
- ▶ Maximizing conditional (log) likelihood

$$P(y_{1:n}|x_{1:n}) = \frac{1}{Z(x_{1:n})} \exp\left(\sum_{t=1}^n s(y_{t-1}, y_t, x_{1:n})\right)$$

$$Z(x_{1:n}) = \sum_{y'} \exp\left(\sum_{t=1}^n s(y'_{t-1}, y'_t, x_{1:n})\right)$$

- ▶ Optimization with gradient descent
  - ▶ The partition function  $Z$  is computed by Forward algorithm
  - ▶ The gradient formula involves expected counts
    - ▶ Can be computed with Forward-Backward
    - ▶ Or we simply let auto-differentiation handle everything (as discussed earlier)



# CRF Supervised Learning

---

- ▶ Learn CRF given annotated sequence  $\{(x_1, y_1), \dots, (x_n, y_n)\}$
- ▶ Minimizing margin-based loss (Structured SVM)

$$L_{SSVM} = \max_{y_{1:n}} (s(y_{1:n}) + \Delta(y_{1:n}, y_{1:n}^*) - s(y_{1:n}^*))$$


- ▶  $\Delta(y, y^*) \geq 0$  is the cost we incur when we predict  $y$  but the truth is  $y^*$
- ▶  $\max_{y_{1:n}} (\dots)$  can be computed with Viterbi if  $\Delta$  is position-wise decomposable, e.g., num of different labels
- ▶ Advantages
  - ▶ take into account the  $\Delta$  cost
  - ▶ focus on the decision boundary instead of the full distribution
- ▶ Optimization --- loss not differentiable
  - ▶ stochastic subgradient descent
  - ▶ quadratic programming (cutting-plane method)



# CRF Unsupervised Learning

---

- ▶ Learn CRF given unannotated sequence  $\{x_1, \dots, x_n\}$
- ▶ Impossible to compute  $P(x_1, \dots, x_n)$  with a CRF!
- ▶ CRF autoencoder (CRF-AE)
  - ▶ Encoder: CRF
  - ▶ Decoder: simply predict each word from its tag





# CRF Unsupervised Learning

---

- ▶ Learn CRF given unannotated sequence  $\{x_1, \dots, x_n\}$
- ▶ Impossible to compute  $P(x_1, \dots, x_n)$  with a CRF!
- ▶ CRF autoencoder (CRF-AE)
  - ▶ Encoder: CRF
  - ▶ Decoder: simply predict each word from its tag
- ▶ Training loss:

$$\begin{aligned} P(\widehat{x}_{1:n} | x_{1:n}) &= \sum_{y_{1:n}} P(y_{1:n} | x_{1:n}) P(\widehat{x}_{1:n} | y_{1:n}) \\ &= \sum_{y_{1:n}} \frac{1}{Z(x_{1:n})} \prod_{t=1}^n \exp(s(y_{t-1}, y_t, x_{1:n})) P(\widehat{x}_t | y_t) \end{aligned}$$

- ▶ The loss can be computed with Forward algorithm and optimized with gradient descent
- 



# CRF in general


---

- ▶ An extension of Markov networks (aka. Markov random fields) where everything is conditioned on the input

$$P(y|x) = \frac{1}{Z(x)} \prod_c \psi_c(y_c, x)$$

- ▶ where  $\psi_c(y_c, x)$  is the potential over clique  $C$  and  $Z(x)$  is the normalization coefficient.

$$Z(x) = \sum_y \prod_c \psi_c(y_c, x)$$






# Neural Sequence Labeling Model



# Simplest neural method

---




- ▶ Predicting labels directly from static word embeddings
  - ▶ Problem 1: it does not utilize the context of each word
  - ▶ Problem 2: it does not utilize relations between neighboring labels



# RNN for sequence labeling


*RNNs in the general sense,  
including LSTMs and GRUs*



- ▶ Predicting labels from RNN hidden vectors
  - ▶ Problem 1: it does not utilize the context of each word
    - ▶ Each hidden vector only incorporates info from the left context
  - ▶ Problem 2: it does not utilize relations between neighboring labels

# Bidirectional RNN

---

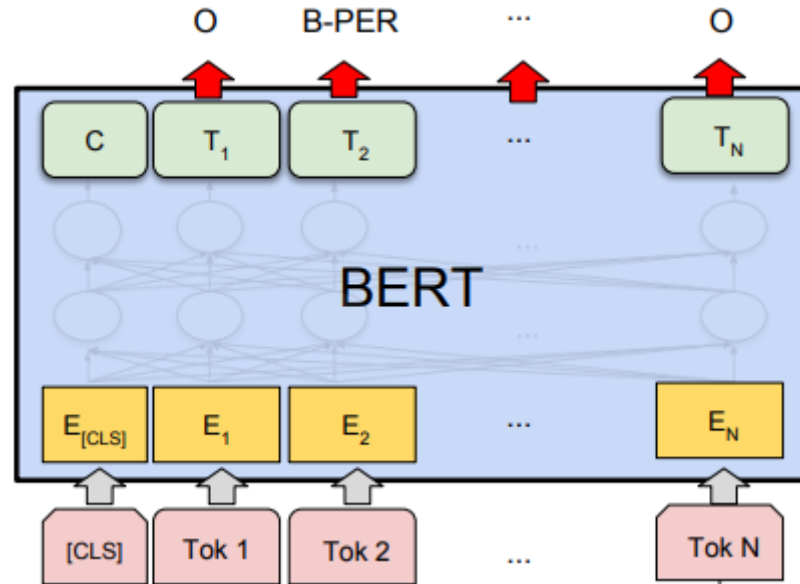


- ▶ Predicting labels from bi-RNN hidden vectors
  - ▶ Problem 1: it does not utilize the context of each word
    - ▶ Solved!
  - ▶ Problem 2: it does not utilize relations between neighboring labels



# Transformer


---



- ▶ Predicting labels from Transformer output vectors
  - ▶ Problem 1: it does not utilize the context of each word
    - ▶ Solved!
  - ▶ Problem 2: it does not utilize relations between neighboring labels

# Neural CRF

---



$$P(y_{1:n}|x_{1:n}) = \frac{1}{Z(x_{1:n})} \prod_{t=1}^n \exp(s(y_{t-1}, y_t, x_{1:n}))$$

- ▶ Use a neural model (RNN, Transformer, or both) to compute CRF potentials (typically only the emission scores)
  - ▶ Both problems solved!
  - ▶ The default model for sequence labeling nowadays





# Inference and Learning

---

- ▶ For *all* these models:
  - ▶ Inference
    - ▶ Without CRF: independent prediction at each position
      - ▶ Sometimes called **neural softmax**
    - ▶ With CRF: Viterbi
  - ▶ Learning
    - ▶ Optimize conditional log likelihood or margin-based loss
    - ▶ Similar to those in CRF learning





# Summary



# Sequence Labeling

---

- ▶ Hidden Markov model (HMM)
  - ▶ Inference: Viterbi, Forward, Backward
  - ▶ Learning: Maximum Likelihood Estimate, Expectation-Maximization / SGD
- ▶ Conditional random field (CRF)
  - ▶ Undirected discriminative models
  - ▶ Inference: Viterbi, Forward, Backward
  - ▶ Learning: conditional likelihood, margin-based loss, CRF-AE
- ▶ Neural models
  - ▶ Neural softmax, neural CRF

